

Election Algorithms

Module 3

Need for a Coordinator

- Many distributed algorithms need one process to act as coordinator – Doesn't matter which process does the job, just need to pick one.
- For example,
 - see the centralized mutual exclusion algorithm.
 - Clock synchronization algorithms
- In general, all processes in the distributed system are equally suitable for the role
- Election algorithms are designed to choose a coordinator.

Election Algorithms

- If we are using one process as a coordinator for a shared resource ...
- ...how do we select that one process?

Solution – an *Election*

- All processes currently involved get together to *choose* a coordinator
- If the coordinator crashes or becomes isolated, elect a new coordinator
- If a previously crashed or isolated process, comes on line, a new election *may* have to be held

Election Algorithms

- Any process can serve as coordinator
- Any process can “call an election”
- (initiate the algorithm to choose a new coordinator).
- There is no harm (other than extra message traffic) in having multiple concurrent elections.
- Elections may be needed when the system is initialized, or if the coordinator crashes or retires.

Assumptions

- Every process/site has a unique ID; e.g.
 - **the network address**
 - **a process number**
- Every process in the system should know the values in the set of ID numbers, although not which processors are up or down.
- The process with the highest ID number will be the new coordinator.

Requirements

- When the election algorithm terminates a single process has been selected and every process knows its identity.
- Formalize: every process p_i has a variable e_i to hold the coordinator's process number.
 - $\forall i, e_i = \text{undefined or } e_i = P$, where P is the non-crashed process with highest id
 - All processes (that have not crashed) eventually set $e_i = P$.

Election Algorithms

- Wired systems
 - Bully algorithm
 - Ring algorithm


Bully Algorithm

- **Key Idea: select process with highest ID**
- Assume
 - All processes know about each other
 - Processes numbered uniquely
 - They do not know each other's state
- Suppose P notices no coordinator
 - Sends *election message* to all higher numbered processes
 - If no response, P takes over as coordinator
 - If any responds, P yields


Bully Algorithm (continued)

- Suppose Q receives *election message*
 - Replies *OK* to sender, saying it will take over
 - Sends a new *election message* to higher numbered processes
- Repeat until only one process left standing
 - Announces victory by sending message saying that it is the coordinator
- Three types of messages
 - Election.
 - OK
 - Coordinator

Bully Algorithm (continued)




(a)




Previous coordinator
has crashed


(b)



(c)



(d)



(e)

Bully Algorithm (continued)

- Suppose R comes back on line
 - Sends a new *election message* to higher numbered processes
- Repeat until only one process left standing
 - Announces victory by sending message saying that it is the coordinator (if not already the coordinator)
- Existing (lower numbered) coordinator yields
 - Hence the term “bully”

Analysis


- Works best if communication in the system has bounded latency so processes can determine that a process has failed by knowing the upper bound (UB) on message transmission time (T) and message processing time (M).
 - $UB = 2 * T + M$
- However, if a process calls an election when the coordinator is still active, the coordinator will win the election.

A Ring Algorithm - Overview

- The ring algorithm assumes that the processes are arranged in a logical ring and each process knows the order of the ring of processes.
- Processes are able to “skip” faulty systems: instead of sending to process j , send to $j + 1$.
- Faulty systems are those that don't respond in a fixed amount of time.


A Ring Algorithm

- P thinks the coordinator has crashed; builds an ELECTION message which contains its own ID number.
- Sends to first live successor
- Each process adds its own number and forwards to next.
- OK to have two elections at once.

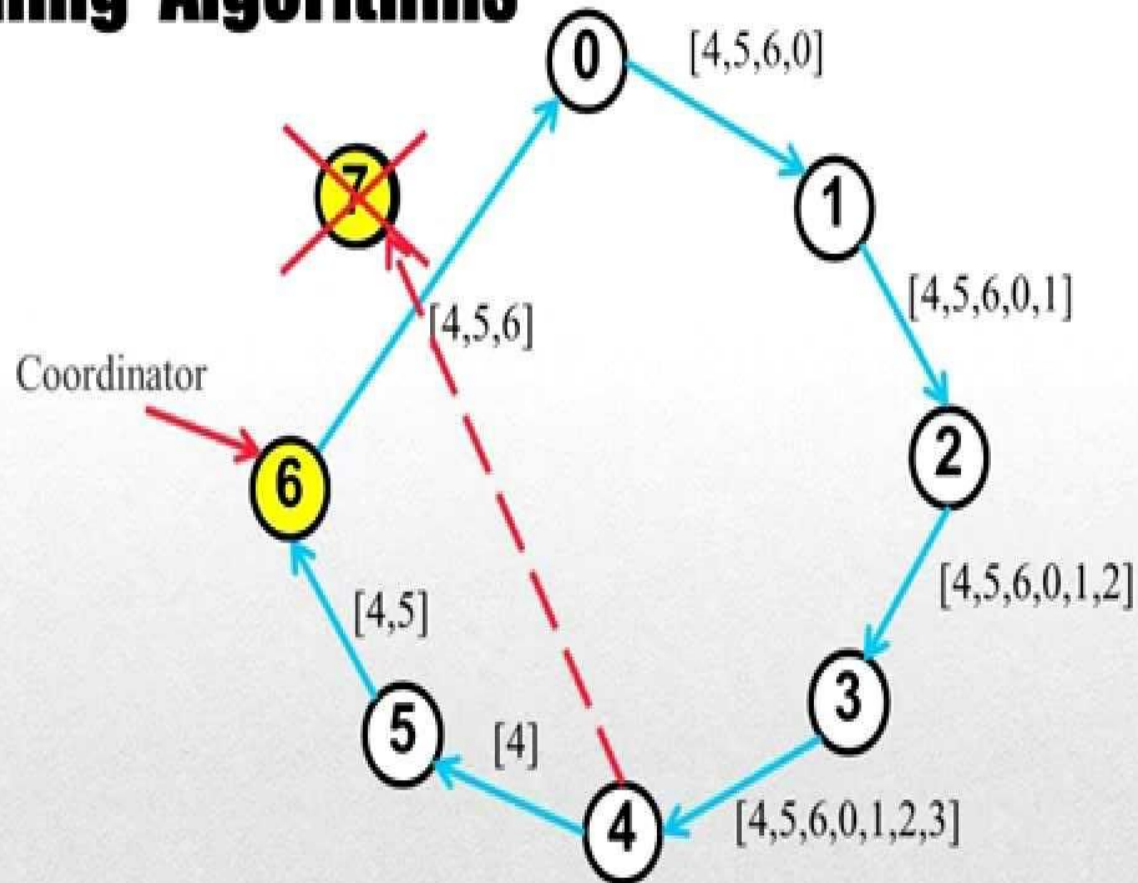


Ring Algorithm - Details

- When the message returns to p , it sees its own process ID in the list and knows that the circuit is complete.
- P circulates a COORDINATOR message with the new high number.
- Here, both 2 and 5 elect 6:
[5,6,0,1,2,3,4]
[2,3,4,5,6,0,1]



Ring Algorithms



Comparison

- Assume n processes and one election in progress

- **Bully Algorithm**

- **Worst case:** initiator is node with lowest ID

Triggers $n-2$ elections at higher ranked nodes:

$O(n^2)$ msgs

- **Best case:** immediate election: $n-2$ messages

- • **Ring Algorithm**

- $2(n-1)$ messages always