

Clock Synchronization in a Distributed System

Topics

- Introduction
- Clock synchronization
- Physical Clock
- Logical clocks
- Election algorithms

Synchronized Distributed Clocks

NEED:

- **Time driven systems:** in statically scheduled systems activities are started at "precise" times in different points of the distributed system.
- **Time stamps:** certain events or messages are associated with a time stamp showing the actual time when they have been produced; certain decisions in the system are based on the "exact" time of the event or event ordering.
- **Calculating the duration of activities:** if such an activity starts on one processor and finishes on another (e.g. transmitting a message), calculating the duration needs clocks to be synchronized.

Lack of Global Time in DS

- It is **impossible** to guarantee that physical clocks run at the same frequency
- Lack of global time, can cause problems
- Example: UNIX make
 - Edit output.c at a client
 - output.o is at a server (compile at server)
 - Client machine clock can be lagging behind the server machine clock

Lack of Global Time – Example



When each machine has its own clock, an event that occurred after another event may nevertheless be assigned an earlier time.

Physical Clock

- Every computer is equipped with CMOS clock circuit. These are electronic devices that count oscillations occurring in a crystal.
- Also called timer, usually a quartz crystal, oscillating at a well defined frequency.
- Timer is associated with two registers:
 - **A Counter and**
 - **a Holding Register,**
- Counter decreasing one at each oscillations. When counter reaches zero, an interrupt is generated; this is the **clock tick**. And again counter is loaded with the value of holding register
- Clock tick have a frequency of 60-100 ticks per second.

Drifting of Clock (cont'd)

- The problems:

1. Crystals cannot be tuned perfectly. Temperature and other external factors can also influence their frequency.



Clock drift: the computer clock differs from the real time.

2. Two crystals are never identical.



Clock skew: the computer clocks on different processors of the distributed system show different time.

Need for Precision Time

- Social networking services
- Stock market buy and sell orders
- Secure document timestamps (with cryptographic certification)
- Aviation traffic control and position reporting
- Radio and TV programming launch and monitoring
- Intruder detection, location and reporting
- Multimedia synchronization for real-time teleconferencing
- Interactive simulation event synchronization and ordering
- Network monitoring, measurement and control
- Early detection of failing network infrastructure devices and air conditioning equipment
- Differentiated services traffic engineering
- Distributed network gaming and training

Clock synchronization




Figure 5-6 Classification of clock synchronization algorithms

Clock Synchronization Algorithms


☞ Distributed Algorithms

- There is no particular time server.
- The processors periodically reach an agreement on the clock value by averaging the time of neighbors clock and its local clock.
 - This can be used if no UTC receiver exists (no external synchronization is needed). Only internal synchronization is performed.
 - Processes can run on different machines and no global clock to judge which event happens first.

Cristian's Algorithm

- *Cristian's Algorithm is centralized algorithm.*
- ❑ The simplest algorithm for setting time, it issues a Remote Procedure Call to time server and obtain the time.
- ❑ A machine sends a request to time server .
- ❑ The time server sends a reply with current UTC when receives the request.
- ❑ The machine measures the time delay between time server sending the message and machine receiving it. Then it uses the measure to adjust the clock.

Centralized algorithm



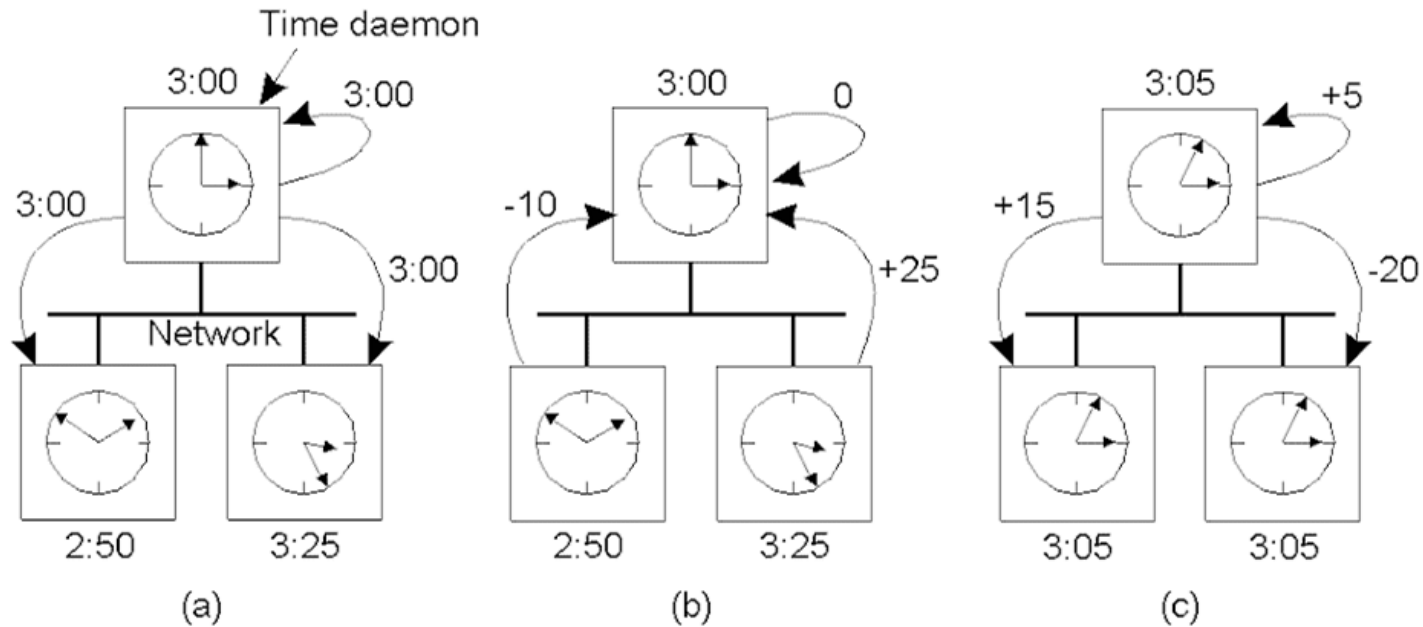
$$\text{Client sets time to: } T_{\text{new}} = T_{\text{server}} + \frac{T_1 - T_0}{2}$$

Figure 5-8 Time approximation using passive time server algorithm

Berkeley Algorithm

- *It is also a Centralized algorithm and its Time server is an Active Machine.*
- ❑ The server polls each machine periodically, asking it for the time.
- ❑ When all the results are in, the master computes the average time.
- ❑ Instead of sending the updated time back to slaves, which would introduce further uncertainty due to network delays, it sends each machine the offset by which its clock needs adjustment.
- ❑ If master machine fails, any other slave could be elected to take over.

Berkeley Algorithm



- a) The time daemon sends synchronization query to other machines in group.
- b) The machines send timestamps as a response to query.
- c) The Server averages the three timestamps and tells everyone how to adjust their clock by sending offsets.

Berkeley's Algorithm

Algorithm

Elect* the master amongst N nodes. Let T_m be the time estimate of the master's clock.

Let $t[i]$ contain the time at each i slave at master

If master


send its T_m along with query for $t[i]$ to slaves; /* for $i = 1 \dots N-1$ */

Adjust = Sum($t[i]$)/ N /* take average including masters


send offset $[i] = \text{Adjust} - t[i]$ to each slave;

If slave


sends query response as $t[i] = T_m - T[i]$; /* for $i = 1 \dots N-1$; calculates the difference between master timestamp T_m , and its own timestamp T^* */



(a) Query-poll



(b) Response



(c) Adjust

Disadvantages of centralized clock

- Subjected to single point failure
- From scalability point of view it is generally not acceptable to get the time requests serviced by single time server.


Network Time Protocol

- NTP is an application layer protocol
- Default port number 123
- Uses standard UDP Internet transport protocol
- Adjust system clock as close to UTC as possible over the Internet.
- Enable sufficiently frequently resynchronizations. (scaling well on large num
- The NTP service is provided by a network of servers located across the Internet.

Network Time Protocol


- Primary servers are connected directly to a time source. (e.g. a radio clock receiving UTC, GPS). bers of clients and servers)
- Secondary servers are synchronized with primary servers.
- The servers are connected in a logical hierarchy called a synchronization subnet. Each level of the synchronization subnet is called stratum.

NTP (Network Time Protocol)-2




- A requests time of B at its own T_1
- B receives request at its T_2 , records
- B responds at its T_3 , sending values of T_2 and T_3
- A receives response at its T_4
- Question: what is $\theta = T_B - T_A$?

NTP (Network Time Protocol)-3




- Question: what is $\theta = T_B - T_A$?
- Assume transit time is approximately the same both ways
- Assume that B is the time server that A wants to synchronize to

NTP (Network Time Protocol)-4



- A knows $(T_4 - T_1)$ from its own clock
- B reports T_3 and T_2 in response to NTP request
- A computes total transit time of

NTP (Network Time Protocol)-5



- One-way transit time is approximately $\frac{1}{2}$ total, i.e.,
- B 's clock at T_4 reads approximately

$$T_3 + \frac{(T_4 - T_1) - (T_3 - T_2)}{2}$$

NTP -6

- Servers organized as *strata*
 - *Stratum 0* server adjusts itself to WWV directly
 - *Stratum 1* adjusts self to *Stratum 0* servers
 - Etc.
- Within a stratum, servers adjust with each other

Summary

- Real synchronization is imperfect.
- Clocks never exactly synchronized.
- Often inadequate for distributed systems
 - might need totally-ordered events
 - might need millionth-of-a-second precision

LOGICAL CLOCKS

Logical Clock

- Assume no central time source –
 - ❑ Each system maintains its own local clock .
 - ❑ Allow to get global ordering on events.
- Assign sequence numbers to messages –
 - ❑ *All cooperating processes can agree on order of event.*

Lamport's Timestamps

- It is used to provide a partial ordering of events with minimal overhead.
- It is used to synchronize the logical clock.
- **It follows some simple rules:**
 - ❖ A process increments its counter before each event in that process i.e. Clock must tick once between every two events.
 - ❖ When a process sends a message, it includes its timestamp with the message.
 - ❖ On receiving a message, the receiver process sets its counter to be the maximum of the message counter and increments its own counter .

'Happened Before' Relation

- a 'Happened Before' b : $a \rightarrow b$

Situations:

1. If a and b are events in the same process, and a comes before b , then $a \rightarrow b$.

1. If

a : event of message sent

b : event of receipt of the same message

then $a \rightarrow b$.

2. **HBR is Transitive:**

If $a \rightarrow b$ and $b \rightarrow c$ then $a \rightarrow c$.

Note: Two distinct events a and b happens in different process that do not exchange messages then these events are said to be *concurrent*

$a \nrightarrow b$ and $b \nrightarrow a$

Implementation of Logical Clocks


Conditions for correct functioning:

- C₁: If a and b are two events in the same process, and $a \rightarrow b$, then we demand that $C(a) < C(b)$.
- C₂: If a corresponds to sending a message m , and b corresponds to receiving that message, then also $C(a) < C(b)$.
- C₃: A clock C associated with the process P must always go forward, never backwards. Hence corrections to a logical clock must be always made by adding a positive value, never subtracting from it.

Lamport logical clocks

- Lamport clock L orders events consistent with logical “happens before” ordering
 - If $e \rightarrow e'$, then $L(e) < L(e')$
- But not the converse
 - $L(e) < L(e')$ does not imply $e \rightarrow e'$
- Similar rules for concurrency
 - $L(e) = L(e')$ implies $e \parallel e'$ (for distinct e, e')
 - $e \parallel e'$ does not imply $L(e) = L(e')$
- i.e., Lamport clocks arbitrarily order some concurrent events

Lamport's Logical Clocks (2)



(a)

Three processes, each with its own clock. The clocks run at different rates.


Lamport's algorithm

- Each process i keeps a local clock, L_i
- Three rules:
 1. At process i , increment L_i before each event
 2. To send a message m at process i , apply rule 1 and then include the current local time in the message: i.e., $send(m, L_i)$
 3. To receive a message (m, t) at process j , set $L_j = \max(L_j, t)$ and then apply rule 1 before time-stamping the receive event
- The global time $L(e)$ of an event e is its local time
 - For an event e at process i , $L(e) = L_i(e)$

Rules for adjusting clocks

- For two events a and b in same process p l
 - $C(b) = C(a) + 1$
- If a is sending process and b is receiving process of p_i and p_j then,
 - $C_j(b) = \max((C_i(a) + 1), C_j(b))$

Lamport's Logical Clocks



Lamport's algorithm corrects the clocks.

Position of logical clocks in Middleware




Figure 5-18 Positioning of Lamport's logical clocks in distributed systems

Vector Timestamps – 1

- In the system of vector clocks, the time domain is represented by a set *non-negative integer vectors*.
- Each process p_i maintains a vector $vt_i[1..n]$, where $vt_i[j]$ is the local logical clock of p_i and describes the logical time progress at process p_i .
- $Vt_i[j]$ represents process p_i 's latest knowledge of process p_j local time. If $vt_i[j] = x$, then process p_i knows that local time at process p_j has progressed till x .
- The entire vector vt_i constitutes p_i 's view of the global logical time and is used to timestamp events.

Vector Timestamps – 2

- Initially, all vectors $[0,0,\dots,0]$
- For event on process i , increment own c_i
- Label message sent with local vector
- When process j receives message with vector $[d_1, d_2, \dots, d_n]$:
 - Set local each local entry k to $\max(c_k, d_k)$
 - Increment value of c_j

Vector Timestamps – 3

- Vector clocks overcome the shortcoming of Lamport logical clocks
 - $L(e) < L(e')$ does not imply e happened before e'
- Vector timestamps are used to timestamp local events
- They are applied in schemes for replication of data

Vector Timestamps – 4

Process p_i uses the following two rules **R1** and **R2** to update its clock:

- **R1: Before executing an event, process p_i updates its local logical time as follows:**

$$vt_i[i] := vt_i[i] + d \quad (d > 0).$$

- **R2 Each message m is piggybacked with the vector clock vt of the sender process at sending time.**

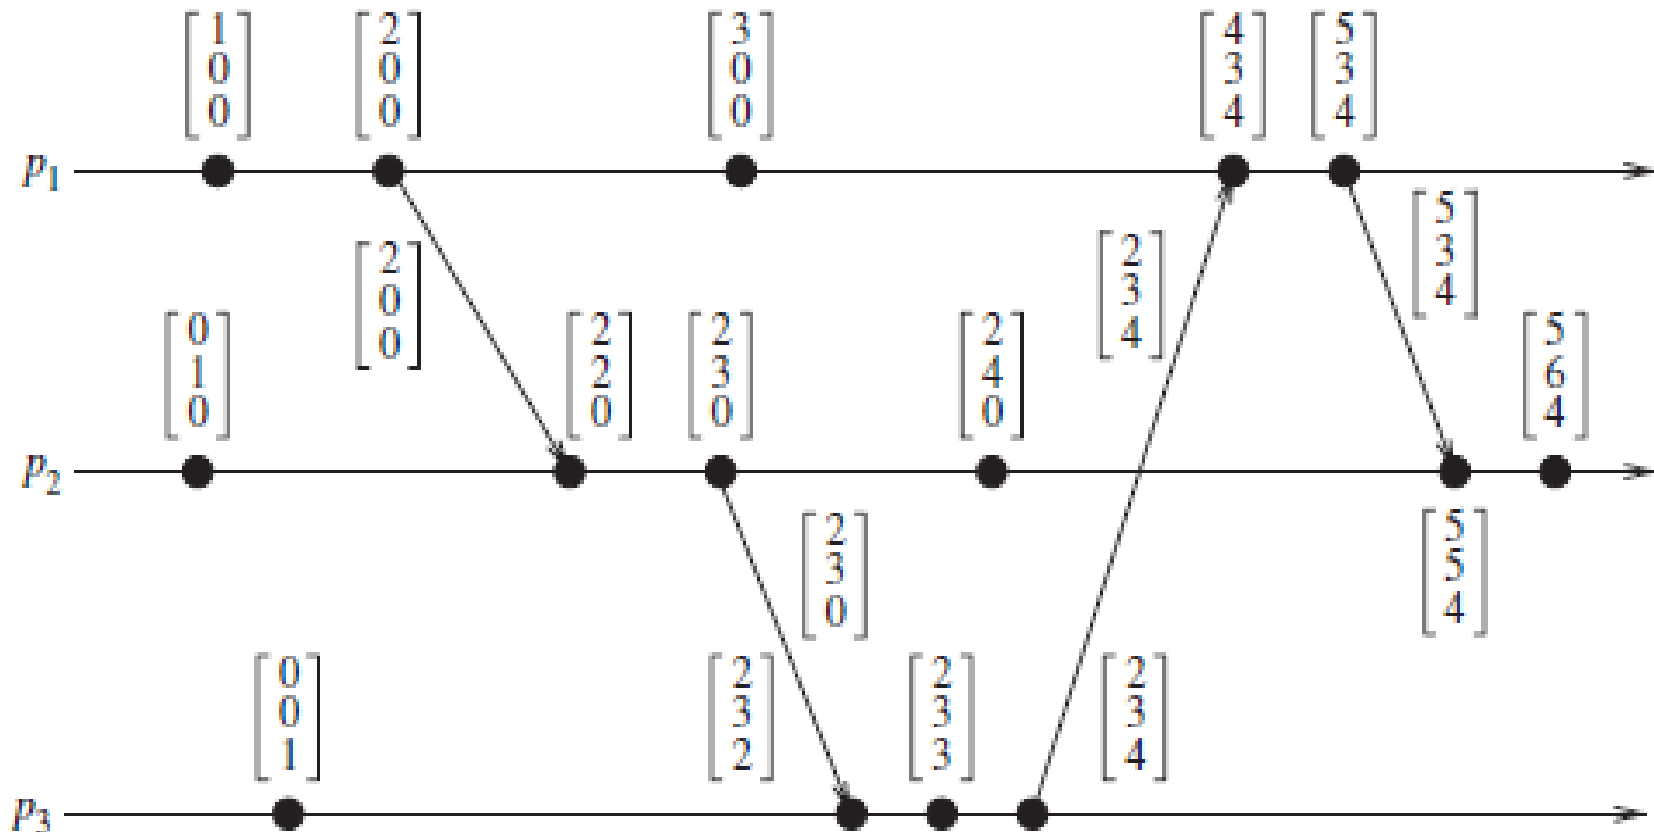
On the receipt of such a message (m, vt) , process p_i executes the following sequence of actions:

1. update its global logical time as follows:

$$1 \leq k \leq n : vt_i[k] := \max(vt_i[k], vt[k]);$$

2. execute **R1**;
3. deliver the message m .

Vector Timestamps – 5(Example)



Vector Timestamps – 6




Figure 5-20 Example of vector timestamp

Important Lessons

- Clocks on different systems will always behave differently
 - Skew and drift between clocks
- Time disagreement between machines can result in undesirable behavior
- Two paths to solution: synchronize clocks or ensure consistent clocks
- Clock synchronization
 - Rely on a time-stamped network messages
 - Estimate delay for message transmission
 - Can synchronize to UTC or to local source