
Module 3

Classification

Supervised vs. Unsupervised Learning

- Supervised learning (classification)
 - Supervision: The training data (observations, measurements, etc.) are accompanied by **labels** indicating the class of the observations
 - New data is classified based on the training set
- Unsupervised learning (clustering)
 - The class labels of training data is unknown
 - Given a set of measurements, observations, etc. with the aim of establishing the existence of classes or clusters in the data

Classification

- **Classification**
 - predicts categorical class labels (discrete or nominal)
 - classifies data (constructs a model) based on the training set and the values (**class labels**) in a classifying attribute and uses it in classifying new data
- Typical applications
 - Credit/loan approval:
 - Medical diagnosis: if a tumor is cancerous or benign
 - Fraud detection: if a transaction is fraudulent
 - Web page categorization: which category it is

Classification: Definition


- Given a collection of records (*training set*)
 - Each record contains a set of *attributes*, one of the attributes is the *class*.
- Find a *model* for class attribute as a function of the values of other attributes.
- Goal: previously unseen records should be assigned a class as accurately as possible.
 - A *test set* is used to determine the accuracy of the model. Usually, the given data set is divided into training and test sets, with training set used to build the model and test set used to validate it.

Classification—A Two-Step Process

- **Model construction**: describing a set of predetermined classes
 - Each tuple/sample is assumed to belong to a predefined class, as determined by the **class label attribute**
 - The set of tuples used for model construction is **training set**
 - The model is represented as classification rules, decision trees, or mathematical formulae
- **Model usage**: for classifying future or unknown objects
 - **Estimate accuracy** of the model
 - The known label of test sample is compared with the classified result from the model
 - **Accuracy** rate is the percentage of test set samples that are correctly classified by the model.
 - **Test set** is independent of training set.
 - If the accuracy is acceptable, use the model to **classify new data**

Process (1): Model Construction

Training
Data



Classification
Algorithms

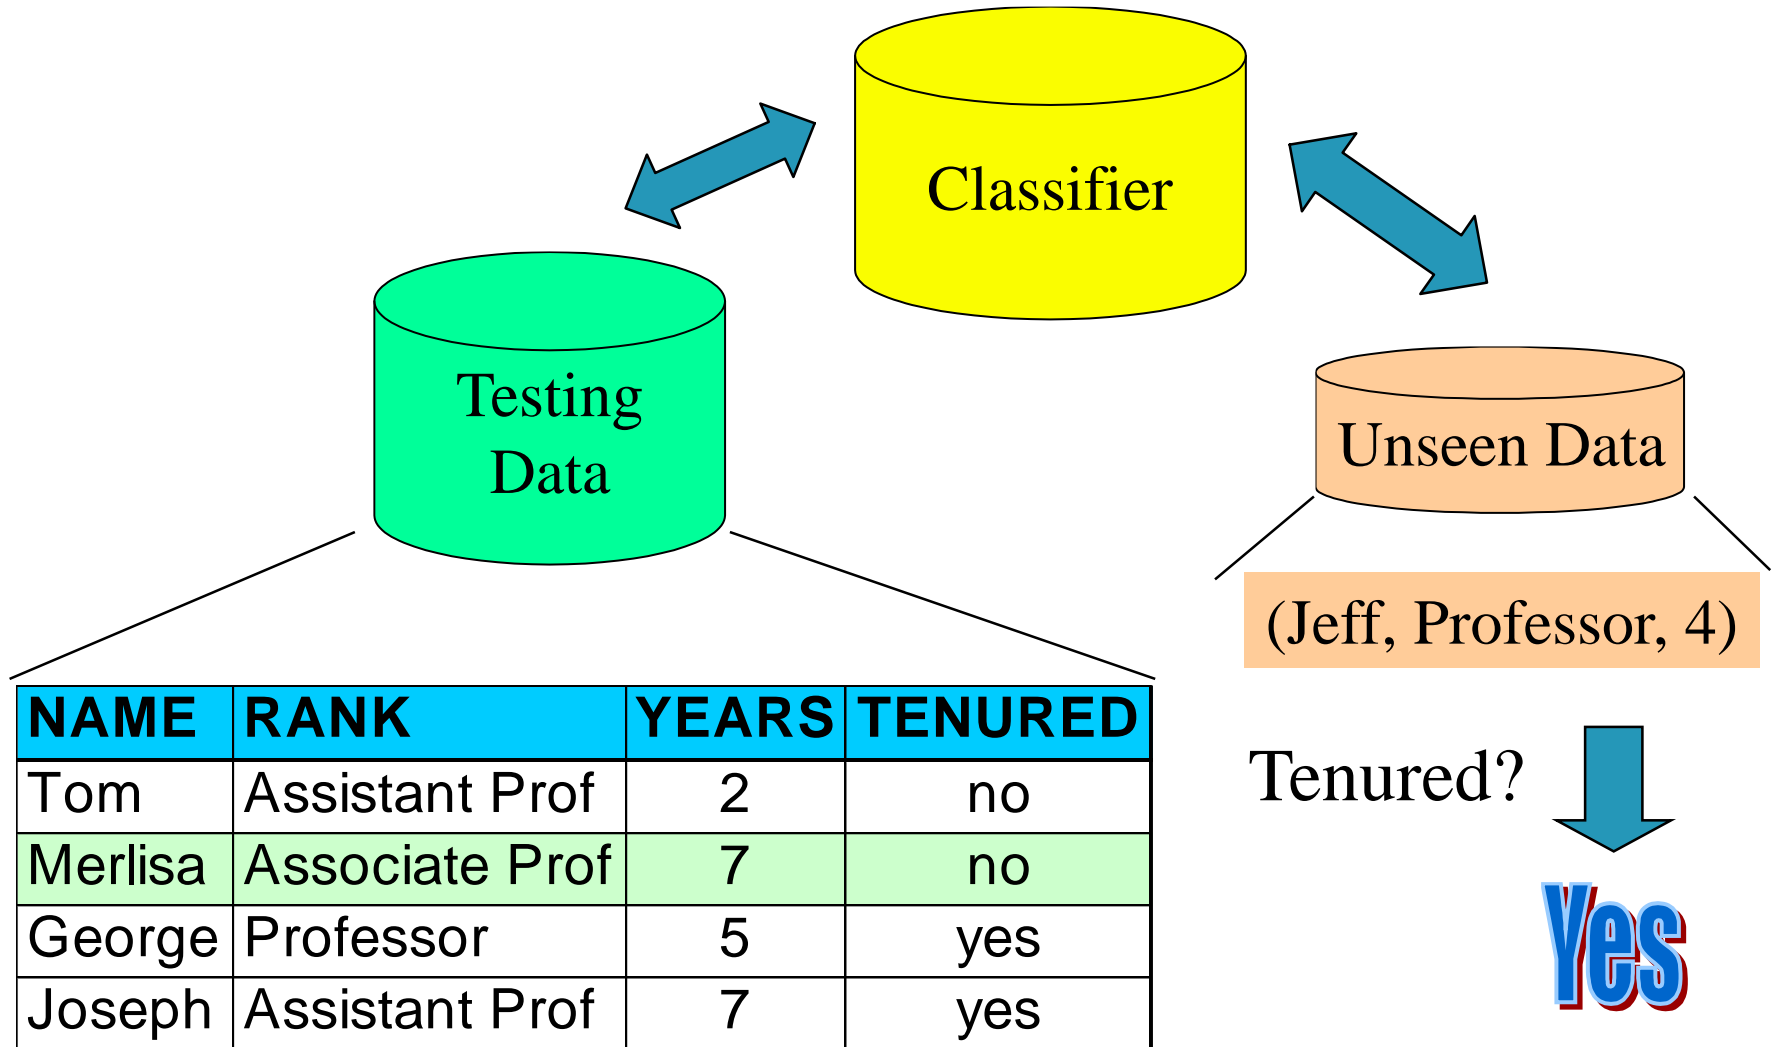


Classifier
(Model)

NAME	RANK	YEARS	TENURED
Mike	Assistant Prof	3	no
Mary	Assistant Prof	7	yes
Bill	Professor	2	yes
Jim	Associate Prof	7	yes
Dave	Assistant Prof	6	no
Anne	Associate Prof	3	no

IF rank = 'professor'
OR years > 6
THEN tenured = 'yes'

Process (2): Using the Model for Classification




Illustrating Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set


Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Examples of Classification

- Predicting tumor cells as benign or malignant
- Classifying credit card transactions as legitimate or fraudulent
- Classifying secondary structures of protein as alpha-helix, beta-sheet, or random coil
- Categorizing news stories as finance, weather, entertainment, sports, etc



Decision Tree Induction

Decision Tree Induction

- J. Ross Quinlan
 - a researcher in machine learning
 - ID3 (Iterative Dichotomiser).
 - Late 1970 and early 1980
- E. B. Hunt, J. Marin , P. T. Stone. And Quinlan
 - presented C4.5 (a successor of ID3),
 - which became a benchmark to which newer supervised learning algorithms are often compared.
- L. Breiman, J. Friedman, R. Olshen, and C. Stone
 - *Classification and Regression Trees* (CART), which described the generation of binary decision trees.

Decision Tree Induction

- ID3, C4.5, and CART adopt a greedy approach in which decision trees are constructed in a top-down recursive divide-and-conquer manner.
- The training set is recursively partitioned into smaller subsets as the tree is being built.

Algorithm : Generate decision tree.

- Generate a decision tree from the training tuples of data partition D .
- Input:
 - Data partition D : which is a set of training tuples and their associated class labels;
 - *attribute list*: the set of candidate attributes;
 - *Attribute selection method*: a procedure to determine the splitting criterion that “best” partitions the data tuples into individual classes. This criterion consists of a *splitting attribute* and, possibly, either a *split point* or *splitting subset*.
- Output: A decision tree.

Algorithm : Generate decision tree.

Method:

- (1) create a node N ;
- (2) if tuples in D are all of the same class, C then
- (3) return N as a leaf node labeled with the class C ;
- (4) if *attribute list* is empty then
- (5) return N as a leaf node labeled with the majority class in D ; // majority voting
- (6) apply Attribute selection method(D , *attribute list*) to find the “best” *splitting criterion*;
- (7) label node N with *splitting criterion*;

Algorithm : Generate decision tree.

- (8) if *splitting attribute* is discrete-valued and
multiway splits allowed then // not restricted to
binary trees
- (9) $attribute\ list \leftarrow attribute\ list - splitting\ attribute;$
// remove *splitting attribute*
- (10) for each outcome j of *splitting criterion*
// partition the tuples and grow subtrees for each partition
- (11) let D_j be the set of data tuples in D satisfying
outcome j ; // a partition
- (12) if D_j is empty then
- (13) attach a leaf labeled with the majority class
in D to node N ;

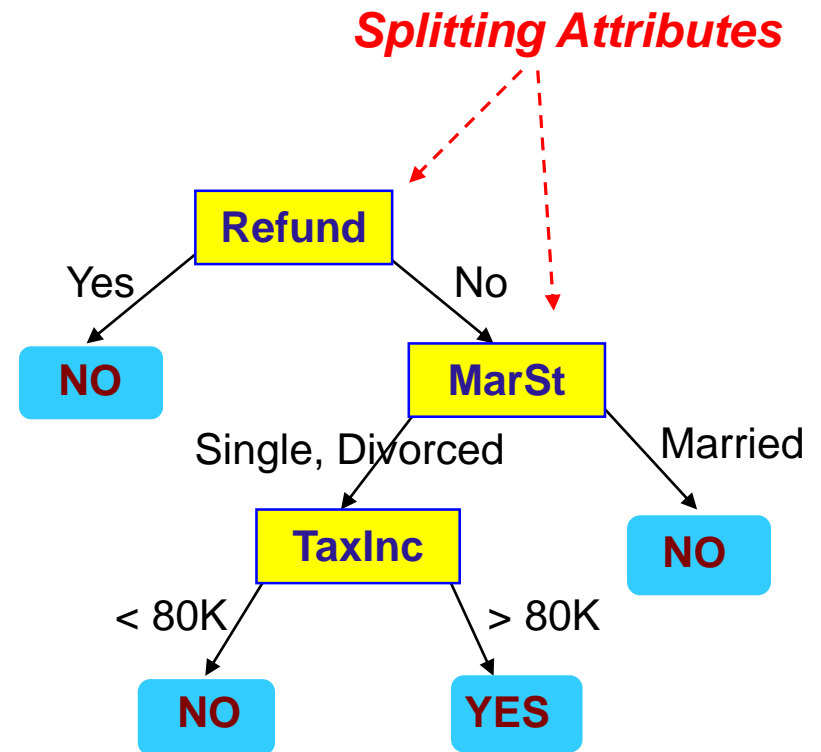
Algorithm : Generate decision tree.

- (14) else attach the node returned by
 Generate _decision _tree(D_j , *attribute list*) to node N ;
 endfor
- (15) return N ;

Example of a Decision Tree

<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Cheat</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes

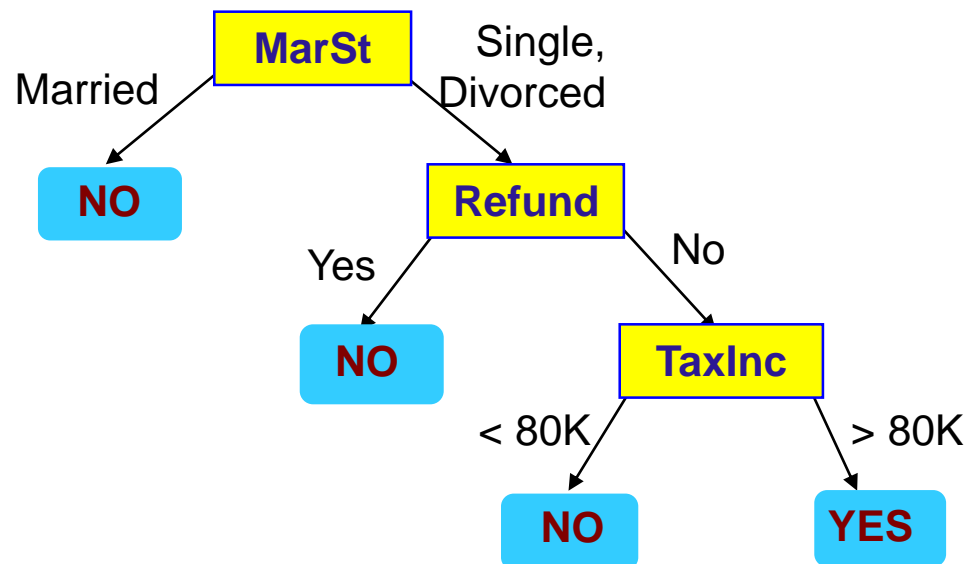
Training Data



Model: Decision Tree

Another Example of Decision Tree

<i>Tid</i>	<i>Refund</i>	<i>Marital Status</i>	<i>Taxable Income</i>	<i>Cheat</i>
1	Yes	Single	125K	No
2	No	Married	100K	No
3	No	Single	70K	No
4	Yes	Married	120K	No
5	No	Divorced	95K	Yes
6	No	Married	60K	No
7	Yes	Divorced	220K	No
8	No	Single	85K	Yes
9	No	Married	75K	No
10	No	Single	90K	Yes



There could be more than one tree that fits the same data!

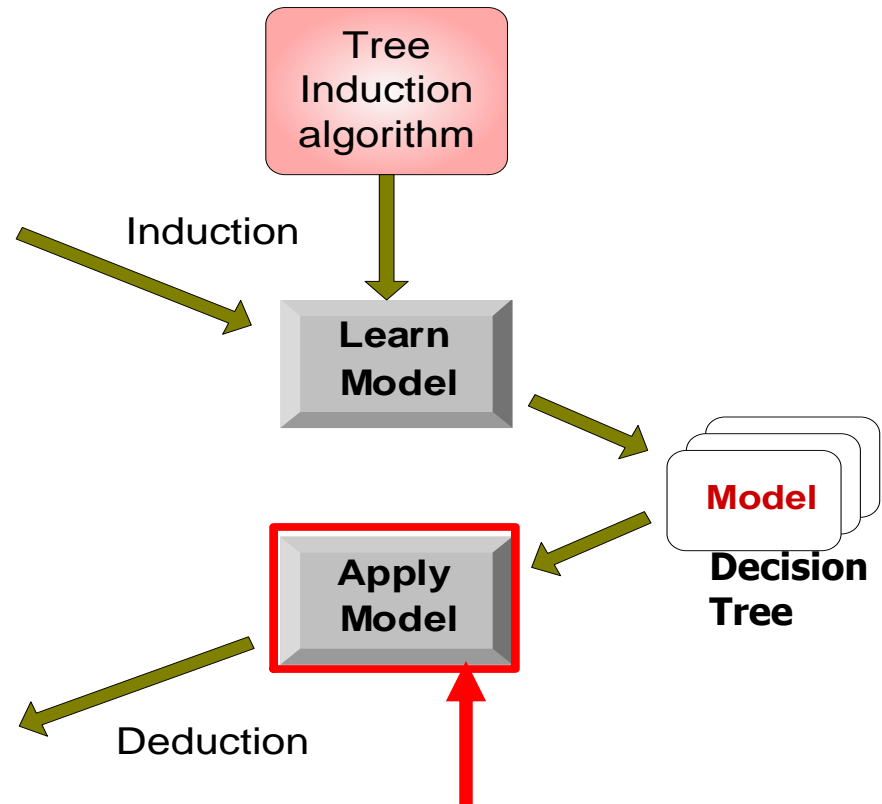
Decision Tree Classification Task

Tid	Attrib1	Attrib2	Attrib3	Class
1	Yes	Large	125K	No
2	No	Medium	100K	No
3	No	Small	70K	No
4	Yes	Medium	120K	No
5	No	Large	95K	Yes
6	No	Medium	60K	No
7	Yes	Large	220K	No
8	No	Small	85K	Yes
9	No	Medium	75K	No
10	No	Small	90K	Yes

Training Set

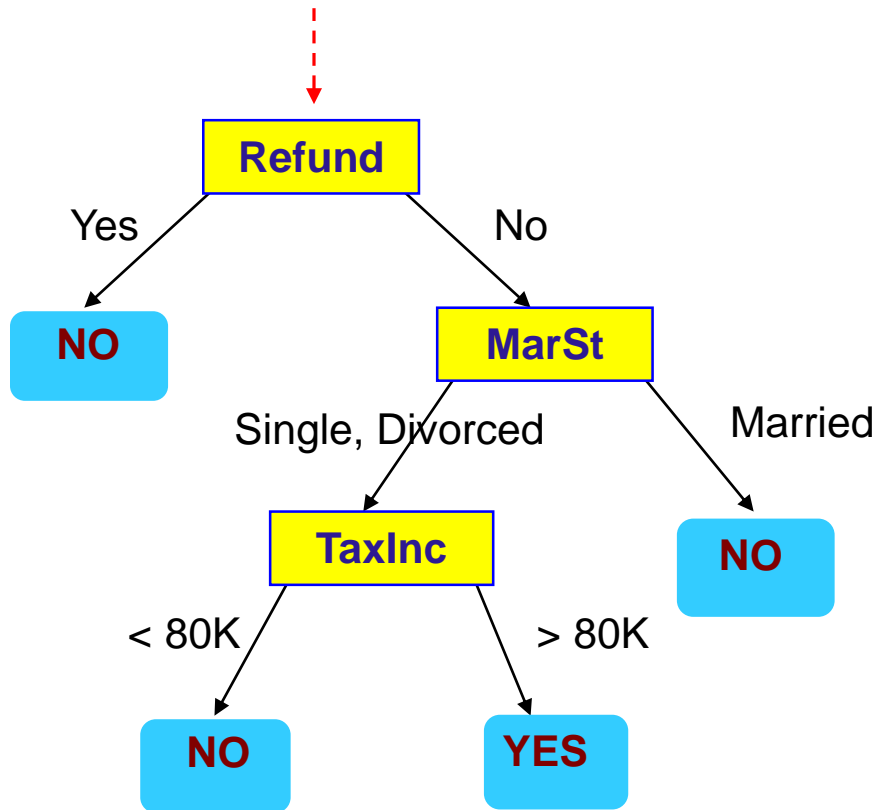
Tid	Attrib1	Attrib2	Attrib3	Class
11	No	Small	55K	?
12	Yes	Medium	80K	?
13	Yes	Large	110K	?
14	No	Small	95K	?
15	No	Large	67K	?

Test Set



Apply Model to Test Data

Start from the root of tree.



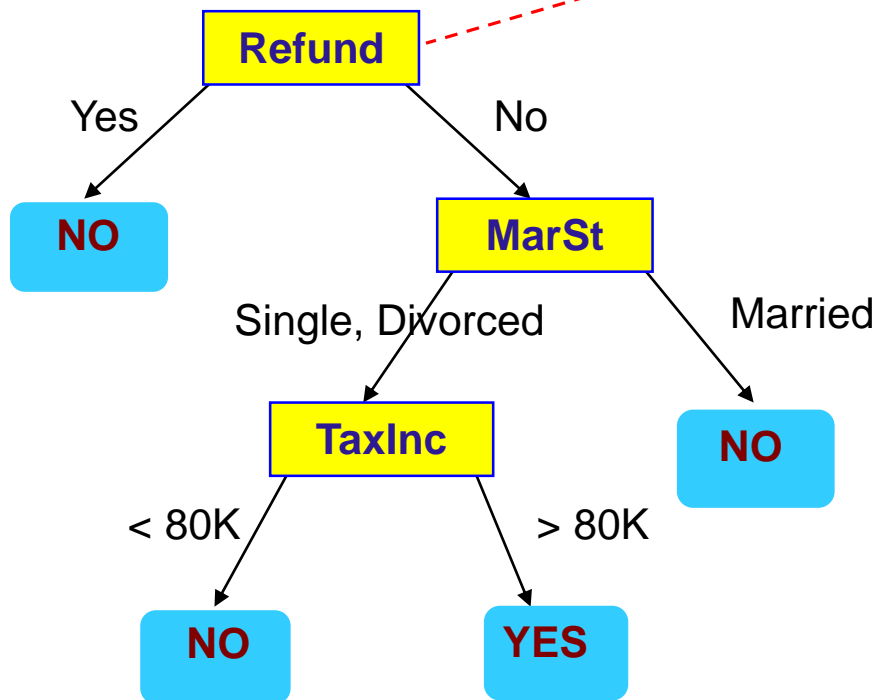
Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?

Apply Model to Test Data

Test Data

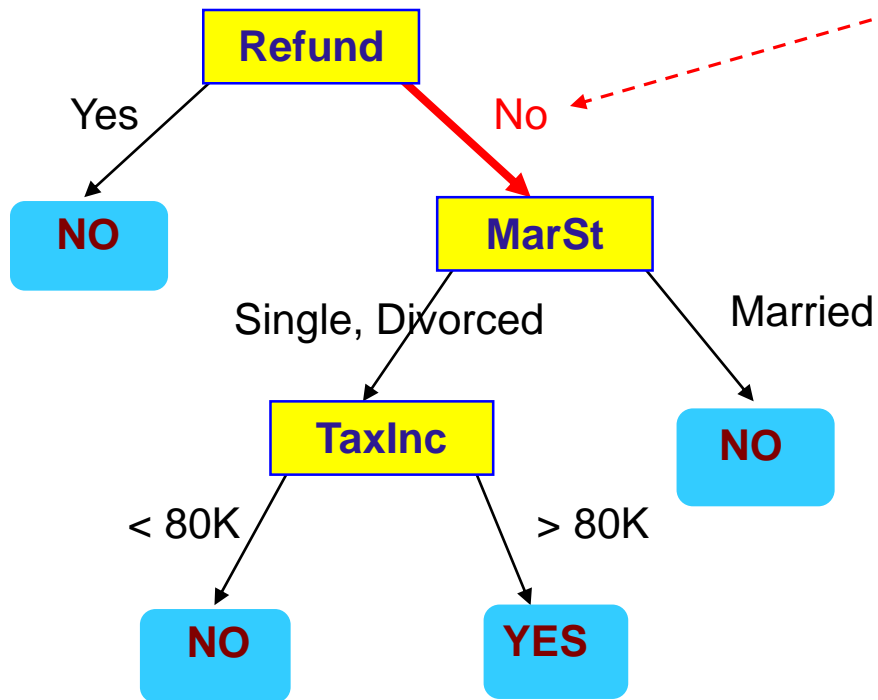
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

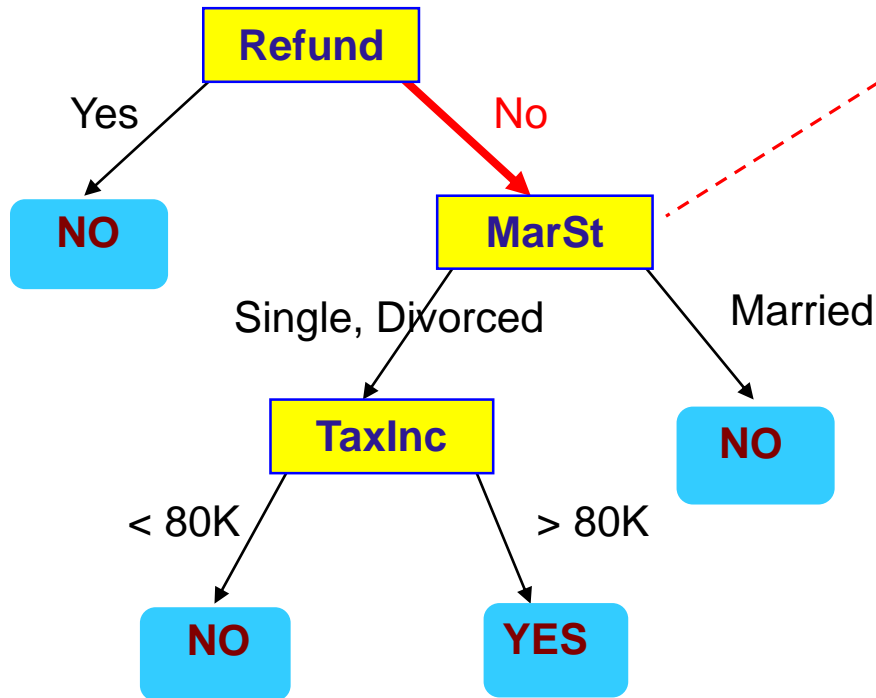
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

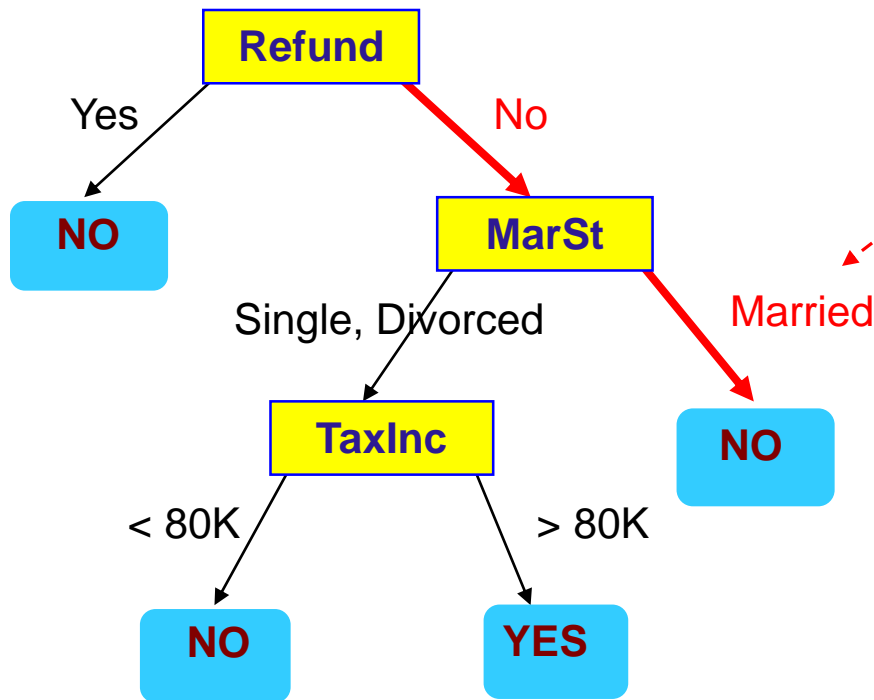
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

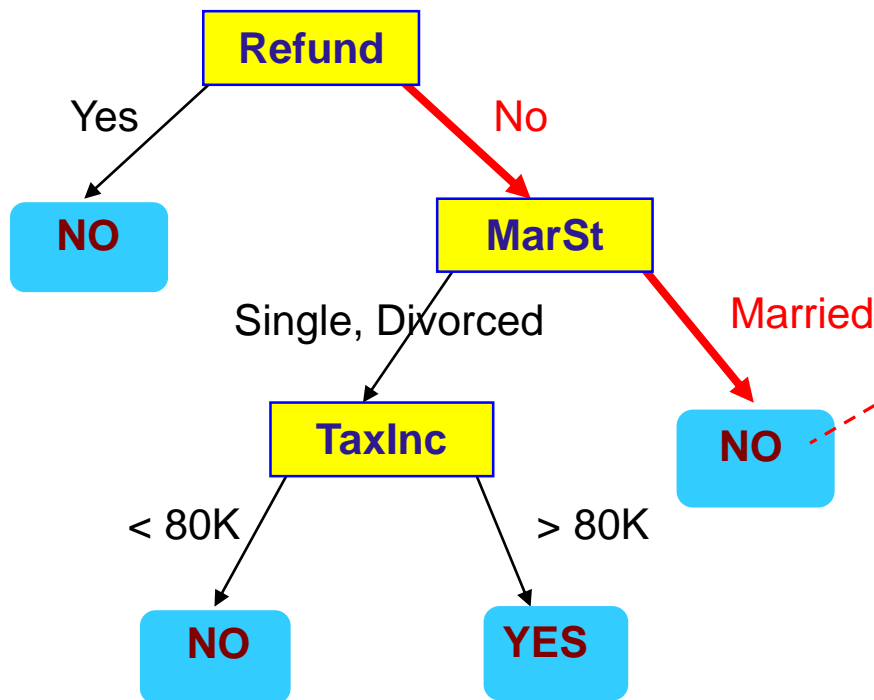
Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Apply Model to Test Data

Test Data

Refund	Marital Status	Taxable Income	Cheat
No	Married	80K	?



Assign Cheat to "No"

Algorithm for Decision Tree Induction

- Basic algorithm (a greedy algorithm)
 - Tree is constructed in a **top-down recursive divide-and-conquer manner**
 - At start, all the training examples are at the root
 - Attributes are categorical (if continuous-valued, they are discretized in advance)
 - Examples are partitioned recursively based on selected attributes
 - Test attributes are selected on the basis of a heuristic or statistical measure (e.g., **information gain**)
- Conditions for stopping partitioning
 - All samples for a given node belong to the same class
 - There are no remaining attributes for further partitioning – **majority voting** is employed for classifying the leaf
 - There are no samples left

Table 1. A small training set

No.	Attributes				Class
	Outlook	Temperature	Humidity	Windy	
1	sunny	hot	high	false	N
2	sunny	hot	high	true	N
3	overcast	hot	high	false	P
4	rain	mild	high	false	P
5	rain	cool	normal	false	P
6	rain	cool	normal	true	N
7	overcast	cool	normal	true	P
8	sunny	mild	high	false	N
9	sunny	cool	normal	false	P
10	rain	mild	normal	false	P
11	sunny	mild	normal	true	P
12	overcast	mild	high	true	P
13	overcast	hot	normal	false	P
14	rain	mild	high	true	N




Figure 2. A simple decision tree

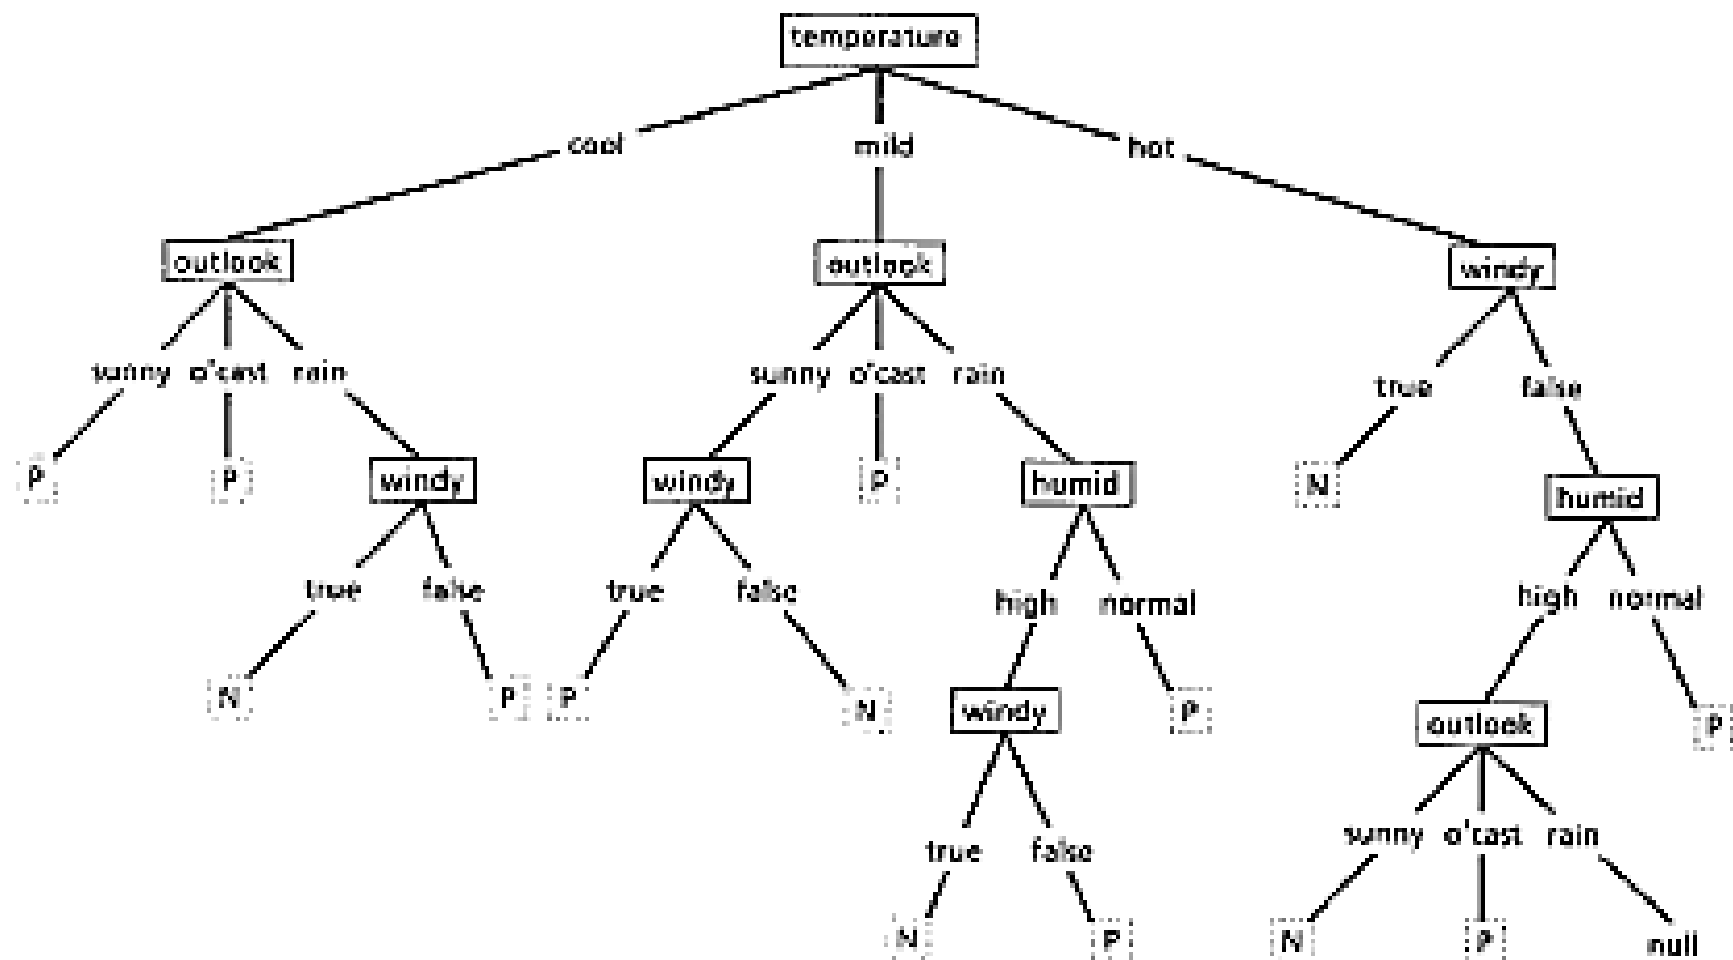


Figure 3. A complex decision tree.

Attribute Selection Measure:

- **Information Gain**
- **Gain Ratio**
- **Gini index**

Attribute Selection Measure: Information Gain

- **ID3** uses **information gain** as its attribute selection measure
- Let Node **N** holds tuples of partition D .
- Select the attribute with the highest information gain as the splitting attributes for Node **N**

Attribute Selection Measure: Information Gain

- **Expected information** (entropy) needed to classify a tuple in D:

$$Info(D) = - \sum_{i=1}^m p_i \log_2(p_i)$$

- Let p_i be the probability that an arbitrary tuple in D belongs to class C_i (distinct class) estimated by $|C_{i,D}|/|D|$
- $C_{i,D}$ be the set of tuples of class C_i in D.
- $|D|$ and $|C_{i,D}|$ denote the number of tuples in D and C_i respectively.
- **$Info(D)$** is also known as **entropy of D**

Attribute Selection Measure: Information Gain (ID3)

- **Information** needed (after using A to split D into v partitions) to classify D:

$$Info_A(D) = \sum_{j=1}^v \frac{|D_j|}{|D|} \times Info(D_j)$$


- **Information gained** by branching on attribute A

$$Gain(A) = Info(D) - Info_A(D)$$


Triangles and Squares

#	Attribute			Shape
	Color	Outline	Dot	
1	green	dashed	no	triange
2	green	dashed	yes	triange
3	yellow	dashed	no	square
4	red	dashed	no	square
5	red	solid	no	square
6	red	solid	yes	triange
7	green	solid	no	square
8	green	dashed	no	triange
9	yellow	solid	yes	square
10	red	solid	no	square
11	green	solid	yes	square
12	yellow	dashed	yes	square
13	yellow	solid	no	square
14	red	dashed	yes	triange

Data Set:
A set of classified objects



Entropy




- 5 triangles
- 9 squares
- class probabilities

- entropy

$$I = -\frac{9}{14} \log_2 \frac{9}{14} - \frac{5}{14} \log_2 \frac{5}{14} = 0.940 \text{ bits}$$

$$I(\text{red}) = -\frac{3}{5} \log_2 \frac{3}{5} - \frac{2}{5} \log_2 \frac{2}{5} = 0.971 \text{ bits}$$

**Entropy
reduction
by
data set
partitioning**



$$I(\text{green}) = -\frac{2}{5} \log_2 \frac{2}{5} - \frac{3}{5} \log_2 \frac{3}{5} = 0.971 \text{ bits}$$

$$I(\text{yellow}) = -\frac{4}{4} \log_2 \frac{4}{4} - \frac{0}{4} \log_2 \frac{0}{4} = 0.0 \text{ bits}$$

$$I = 0.940$$

$$I(\text{red}) = 0.971 \text{ bits}$$

Color?


red

green


yellow

$$I(\text{green}) = 0.971 \text{ bits}$$

$$I(\text{yellow}) = 0.0 \text{ bits}$$




Information Gain



$$\text{Gain}(\text{Color}) = I - I_{\text{res}}(\text{Color}) = 0.940 - 0.694 = 0.246 \text{ bits}$$

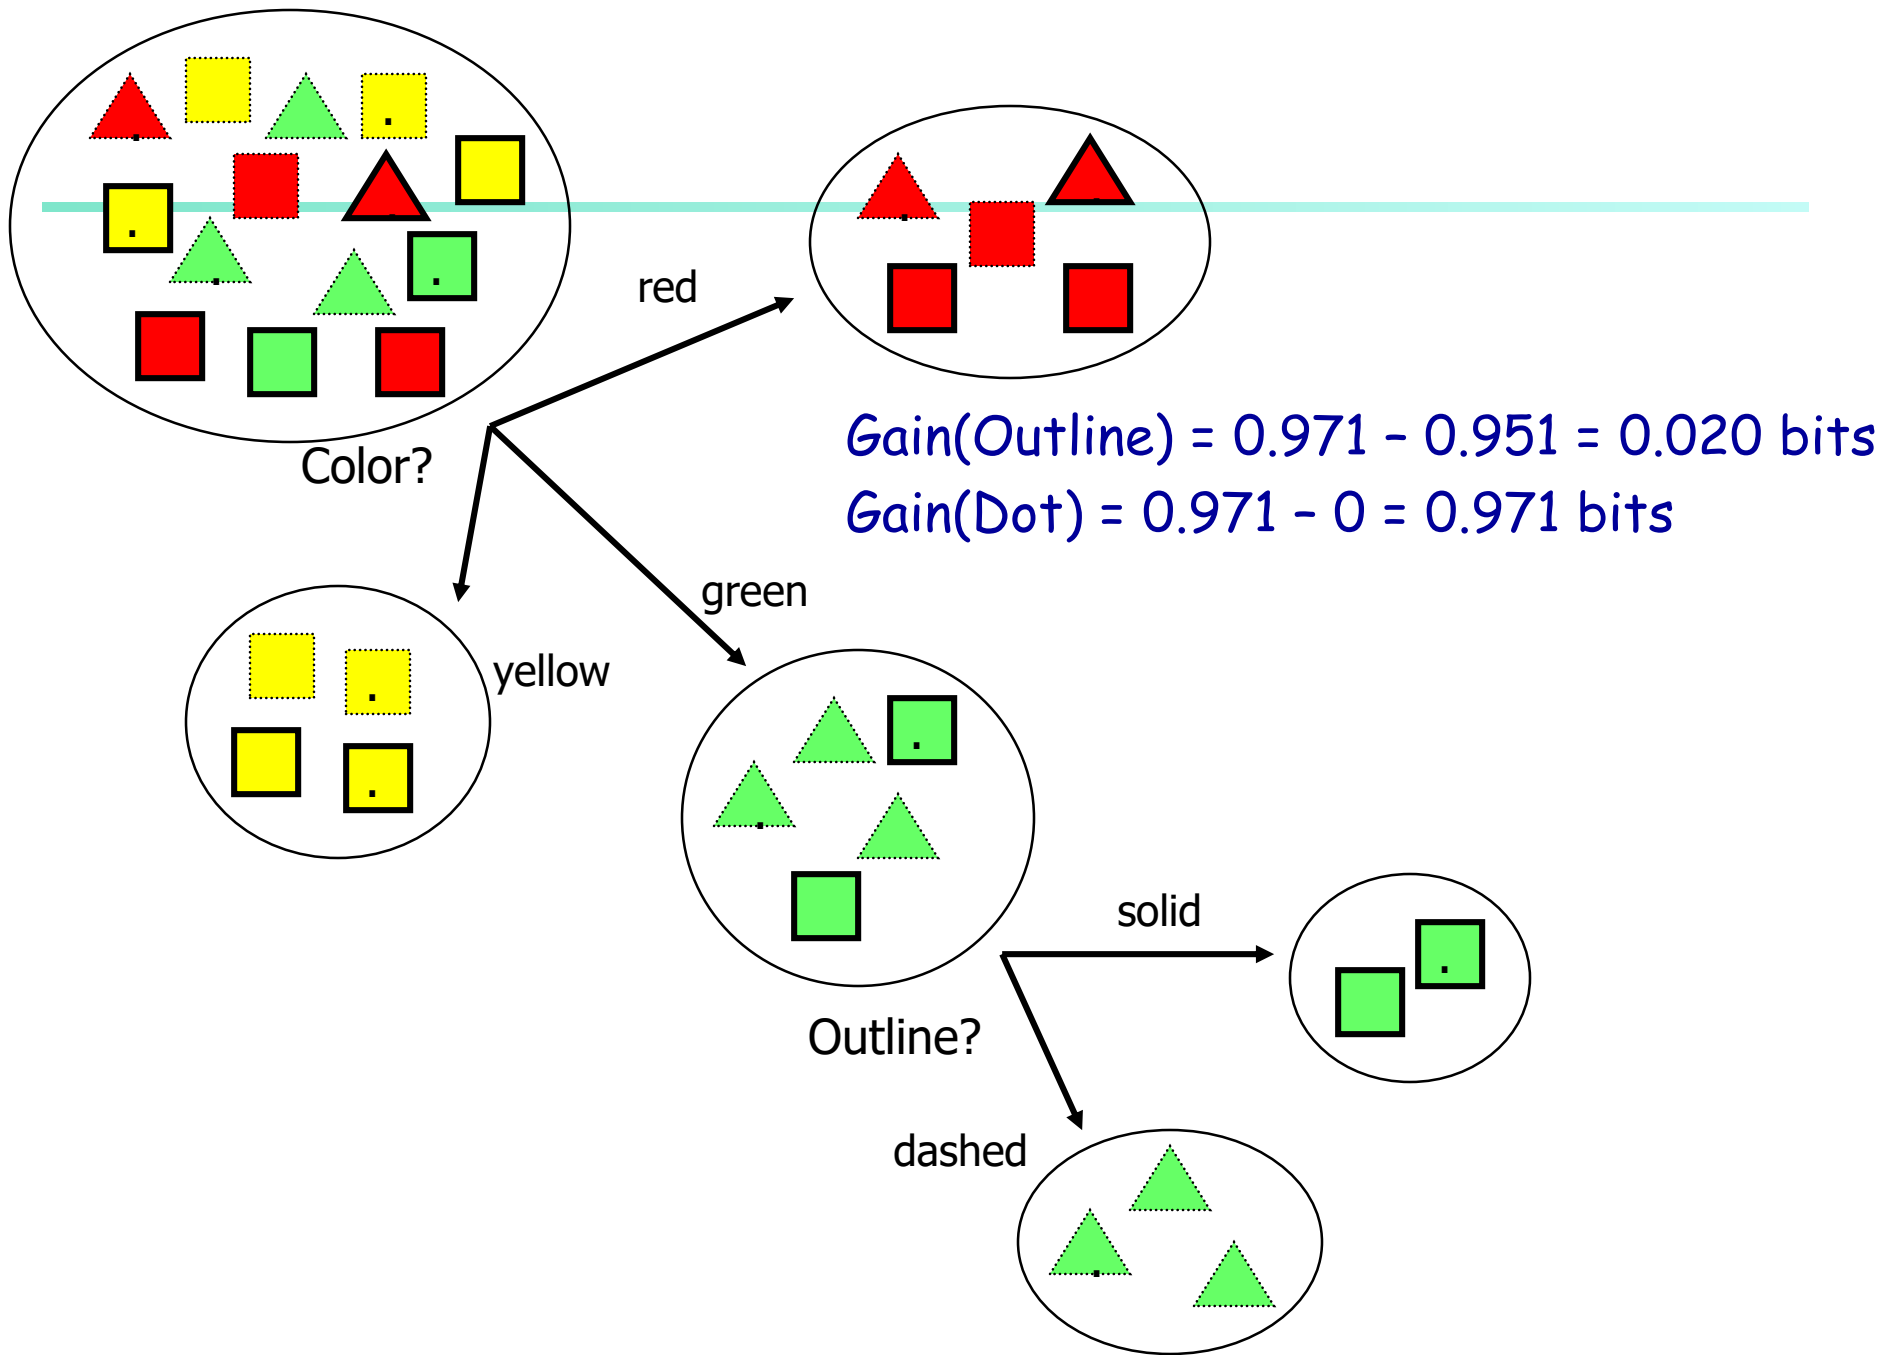
Information Gain of The Attribute

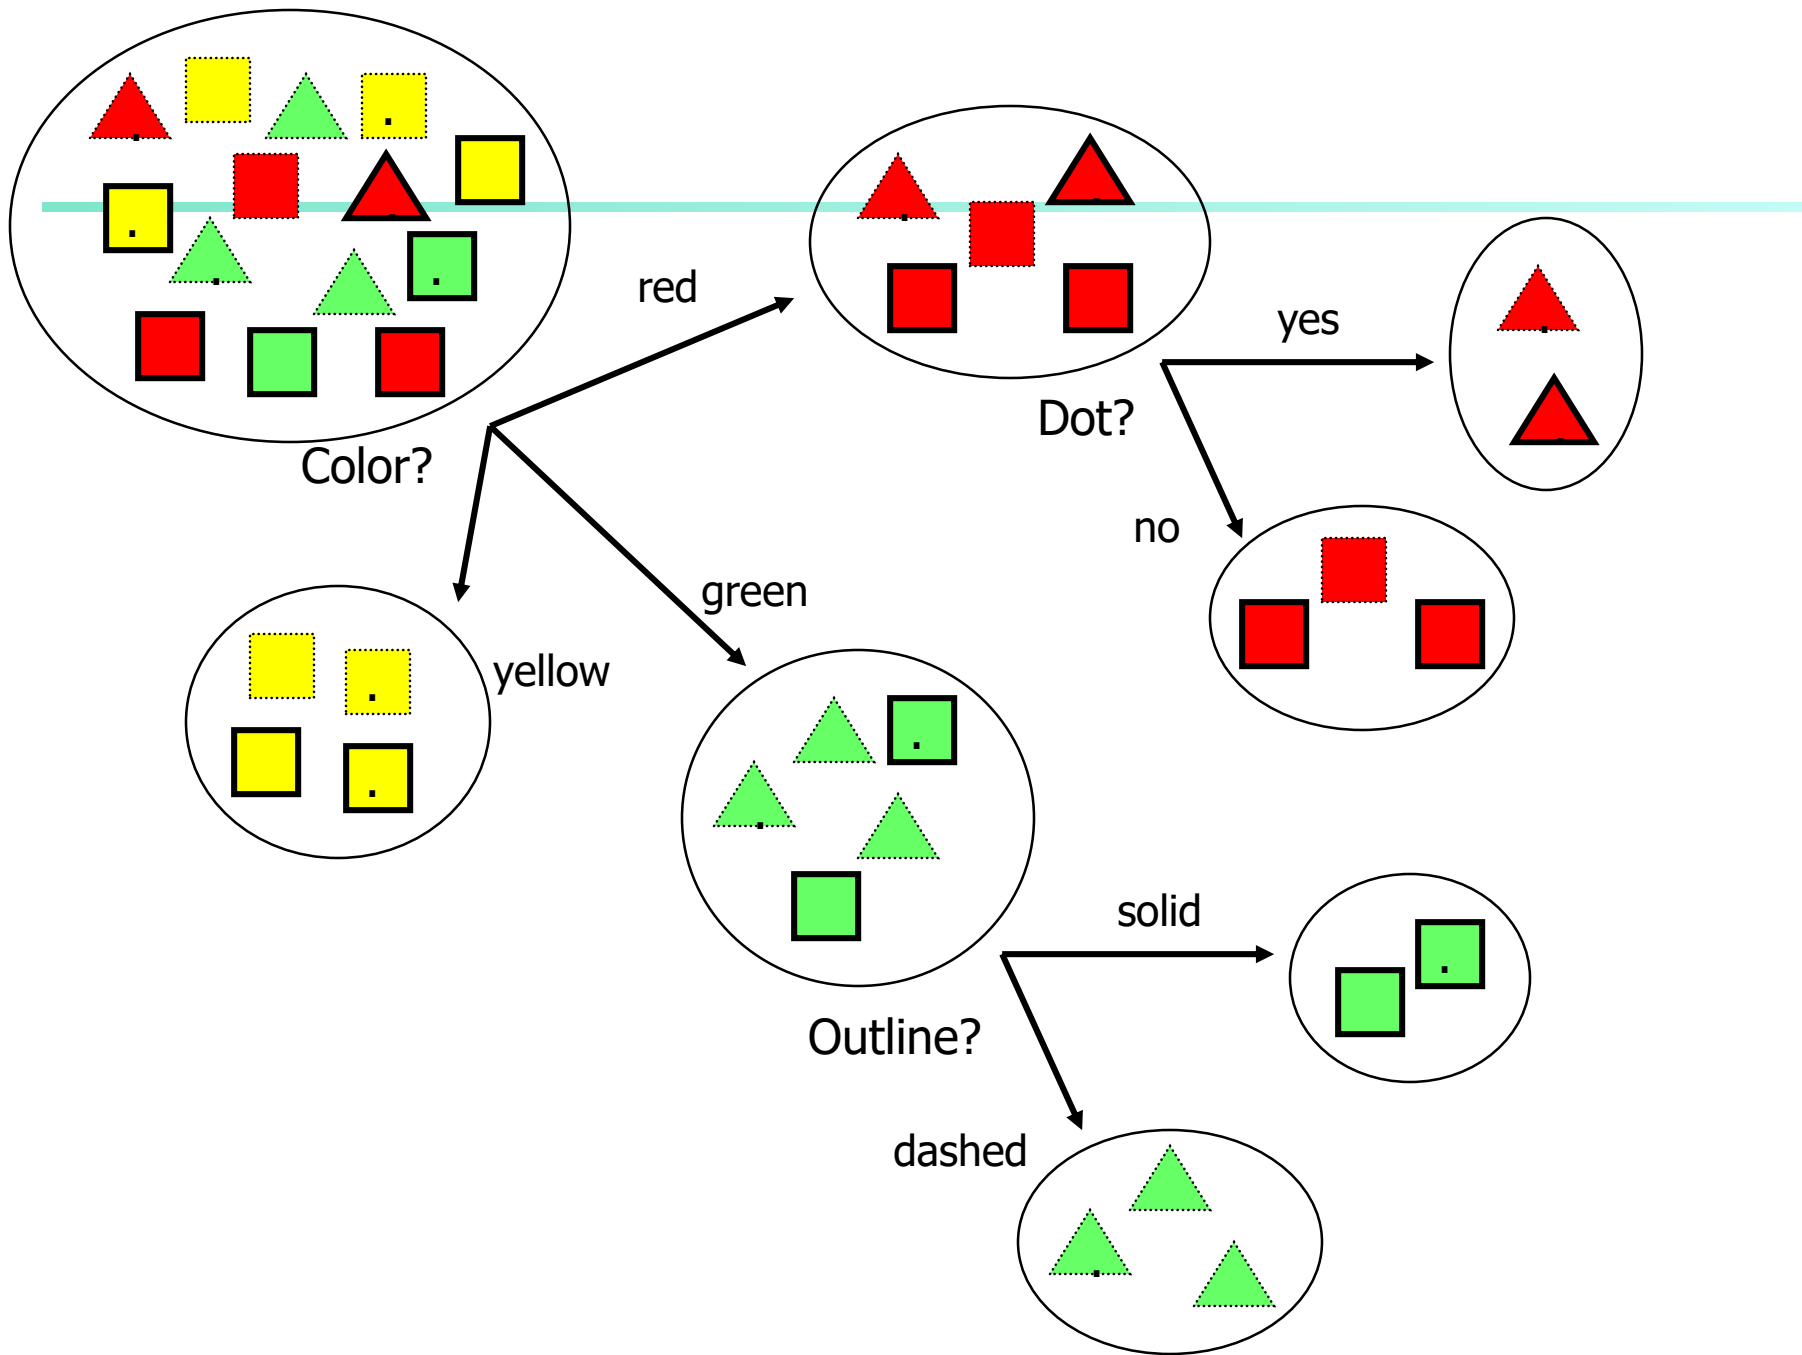
- Attributes
 - $\text{Gain}(\text{Color}) = 0.246$
 - $\text{Gain}(\text{Outline}) = 0.151$
 - $\text{Gain}(\text{Dot}) = 0.048$
- Heuristics: attribute with the highest gain is chosen
- This heuristics is local (local minimization of impurity)



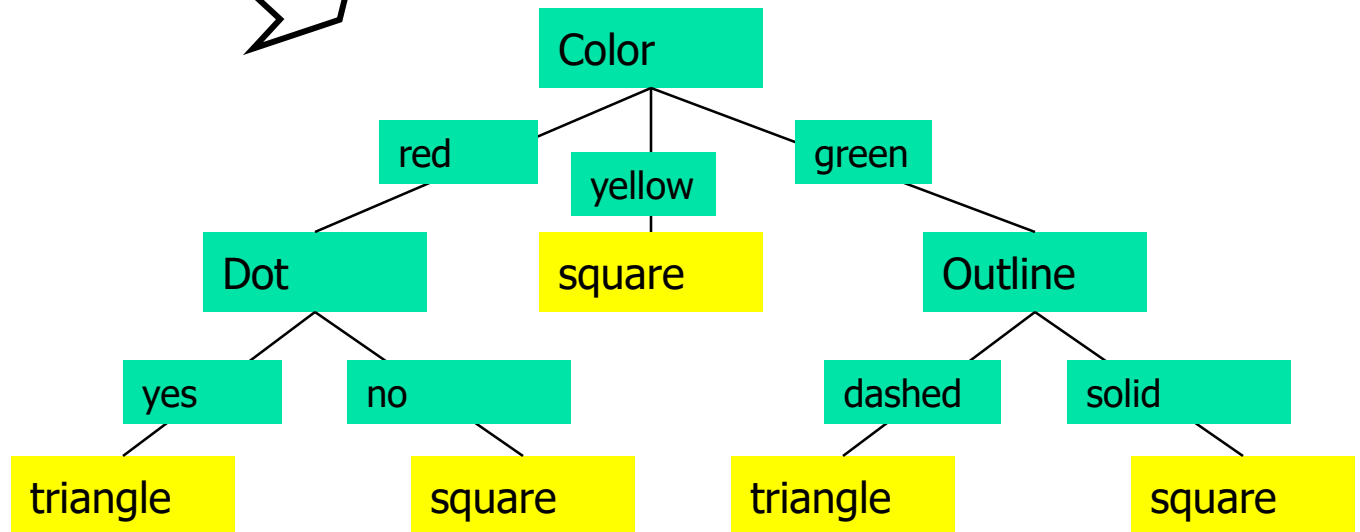


$$\text{Gain(Outline)} = 0.971 - 0 = 0.971 \text{ bits}$$

$$\text{Gain(Dot)} = 0.971 - 0.951 = 0.020 \text{ bits}$$






Decision Tree



Steps to solve ID3 Sum

- compute the entropy for data-set
- for every attribute/feature:
 - 1.calculate entropy for all categorical values
 - 2.take average information entropy for the current attribute
 - 3.calculate gain for the current attribute
- pick the highest gain attribute.
- Repeat until we get the tree we desired.



$$E(\text{Outlook}=\text{sunny}) = -\frac{2}{5} \log\left(\frac{2}{5}\right) - \frac{3}{5} \log\left(\frac{3}{5}\right) = 0.971$$

$$E(\text{Outlook}=\text{overcast}) = -1 \log(1) - 0 \log(0) = 0$$

$$E(\text{Outlook}=\text{rainy}) = -\frac{3}{5} \log\left(\frac{3}{5}\right) - \frac{2}{5} \log\left(\frac{2}{5}\right) = 0.971$$

$$\left. \begin{array}{l} E(\text{Outlook}=\text{sunny}) \\ E(\text{Outlook}=\text{overcast}) \\ E(\text{Outlook}=\text{rainy}) \end{array} \right\} H(S, \text{Outlook})$$


Average Entropy information for Outlook

$$I(\text{Outlook}) = \frac{5}{14} * 0.971 + \frac{4}{14} * 0 + \frac{5}{14} * 0.971 = 0.693$$

$$\text{Gain}(\text{Outlook}) = E(S) - I(\text{outlook}) = 0.94 - 0.693 = 0.247$$

$$\left. \begin{array}{l} I(\text{Outlook}) \\ \text{Gain}(\text{Outlook}) \end{array} \right\} \sum_{t \in T} p(t) H(t)$$

$$\Rightarrow IG(A, S) = H(S) - \sum_{t \in T} p(t) H(t)$$



$$E(\text{Windy}=\text{false}) = -\frac{6}{8} \log\left(\frac{6}{8}\right) - \frac{2}{8} \log\left(\frac{2}{8}\right) = 0.811$$

$$E(\text{Windy}=\text{true}) = -\frac{3}{6} \log\left(\frac{3}{6}\right) - \frac{3}{6} \log\left(\frac{3}{6}\right) = 1$$

Average entropy information for Windy

$$I(\text{Windy}) = \frac{8}{14} * 0.811 + \frac{6}{14} * 1 = 0.892$$

$$\text{Gain}(\text{Windy}) = E(S) - I(\text{Windy}) = 0.94 - 0.892 = 0.048$$

Attribute Selection: Information Gain

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no

Attribute Selection: Information Gain

- Class P: buys_computer = “yes”
- Class N: buys_computer = “no”

$$Info(D) = I(9,5) = -\frac{9}{14} \log_2\left(\frac{9}{14}\right) - \frac{5}{14} \log_2\left(\frac{5}{14}\right) = 0.940$$

age	p _i	n _i	I(p _i , n _i)
<=30	2	3	0.971
31...40	4	0	0
>40	3	2	0.971

$$Info_{age}(D) = \frac{5}{14} I(2,3) + \frac{4}{14} I(4,0) + \frac{5}{14} I(3,2) = 0.694$$

$\frac{5}{14} I(2,3)$ means “age <=30” has 5 out of 14 samples, with 2 yes’es and 3 no’s. Hence

$$Gain(age) = Info(D) - Info_{age}(D) = 0.246$$

Similarly,

$$Gain(income) = 0.029$$

$$Gain(student) = 0.151$$

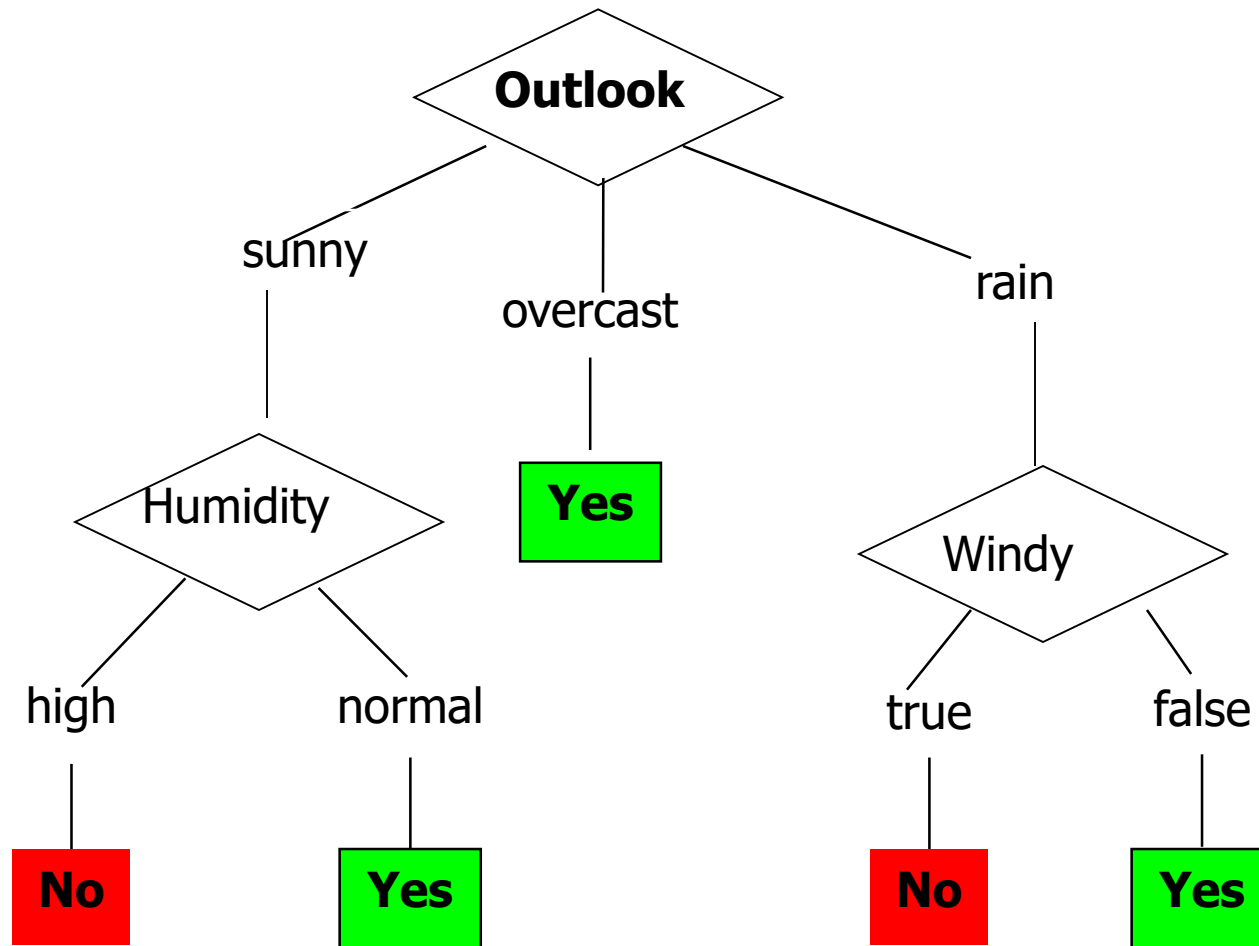
$$Gain(credit_rating) = 0.048$$

age	income	student	credit_rating	buys_computer
<=30	high	no	fair	no
<=30	high	no	excellent	no
31...40	high	no	fair	yes
>40	medium	no	fair	yes
>40	low	yes	fair	yes
>40	low	yes	excellent	no
31...40	low	yes	excellent	yes
<=30	medium	no	fair	no
<=30	low	yes	fair	yes
>40	medium	yes	fair	yes
<=30	medium	yes	excellent	yes
31...40	medium	no	excellent	yes
31...40	high	yes	fair	yes
>40	medium	no	excellent	no





Weather Data: Play or not Play?

Outlook	Temperature	Humidity	Windy	Play?
sunny	hot	high	false	No
sunny	hot	high	true	No
overcast	hot	high	false	Yes
rain	mild	high	false	Yes
rain	cool	normal	false	Yes
rain	cool	normal	true	No
overcast	cool	normal	true	Yes
sunny	mild	high	false	No
sunny	cool	normal	false	Yes
rain	mild	normal	false	Yes
sunny	mild	normal	true	Yes
overcast	mild	high	true	Yes
overcast	hot	normal	false	Yes
rain	mild	high	true	No

Example Tree for “Play?”



Which attribute to select?



Example: attribute “Outlook”


- “Outlook” = “Sunny”:

$$\text{info}([2,3]) = \text{entropy}(2/5, 3/5) = -2/5 \log(2/5) - 3/5 \log(3/5) = 0.971 \text{ bits}$$

- “Outlook” = “Overcast”:

$$\text{info}([4,0]) = \text{entropy}(1) = -1 \log(1) - 0 \log(0) = 0 \text{ bits}$$

Note: $\log(0)$ is not defined, but we evaluate $0 \cdot \log(0)$ as zero



- “Outlook” = “Rainy”:

$$\text{info}([3,2]) = \text{entropy}(3/5, 2/5) = -3/5 \log(3/5) - 2/5 \log(2/5) = 0.971 \text{ bits}$$

- Expected information for attribute:

$$\begin{aligned} \text{info}([3,2],[4,0],[3,2]) &= (5/14) \times 0.971 + (4/14) \times 0 + (5/14) \times 0.971 \\ &= 0.693 \text{ bits} \end{aligned}$$

Computing the information gain

- Information gain:

(information before split) – (information after split)

$$\text{gain("Outlook")} = \text{info}([9,5]) - \text{info}([2,3],[4,0],[3,2]) = 0.940 - 0.693 \\ = 0.247 \text{ bits}$$

- Information gain for attributes from weather data:


$$\text{gain("Outlook")} = 0.247 \text{ bits}$$

$$\text{gain("Temperature")} = 0.029 \text{ bits}$$

$$\text{gain("Humidity")} = 0.152 \text{ bits}$$

$$\text{gain("Windy")} = 0.048 \text{ bits}$$

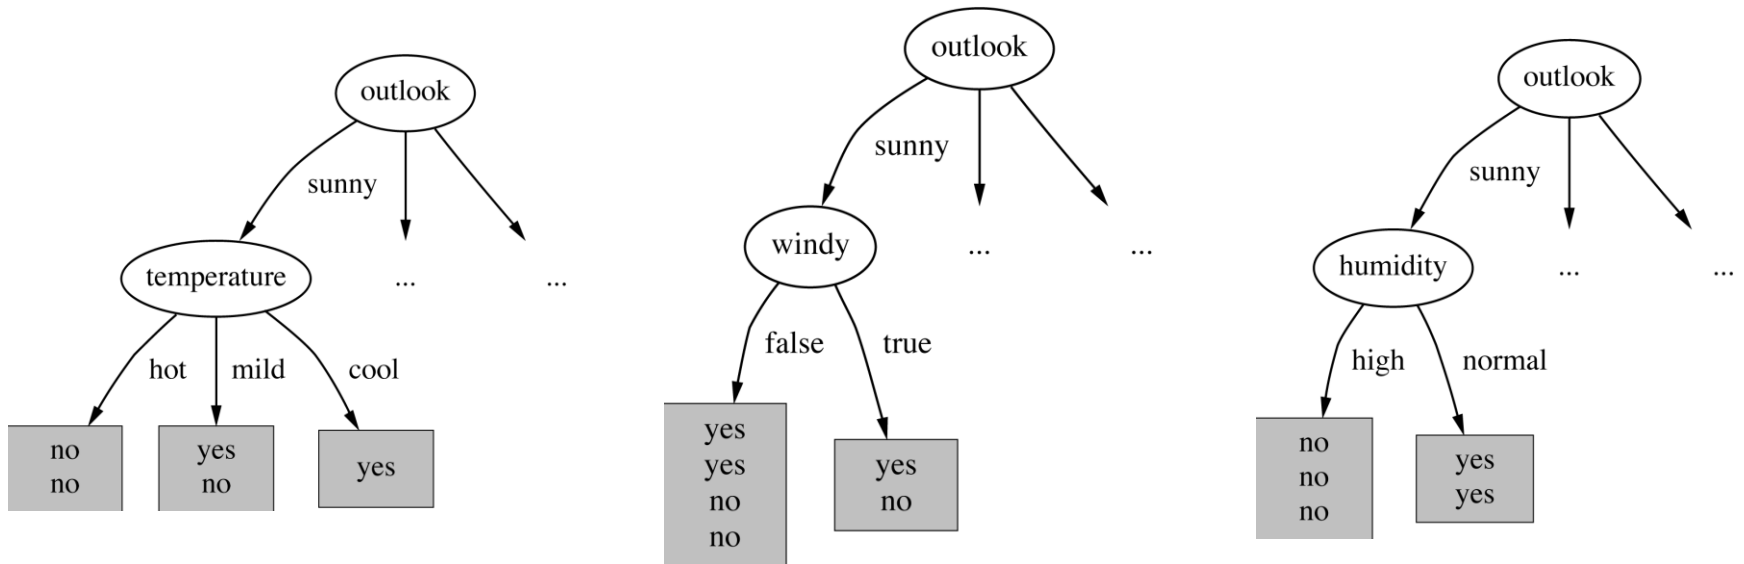
Outlook is selected
as the root node



further splitting
necessary

Outlook = overcast
contains only
examples of class yes

Continuing to split

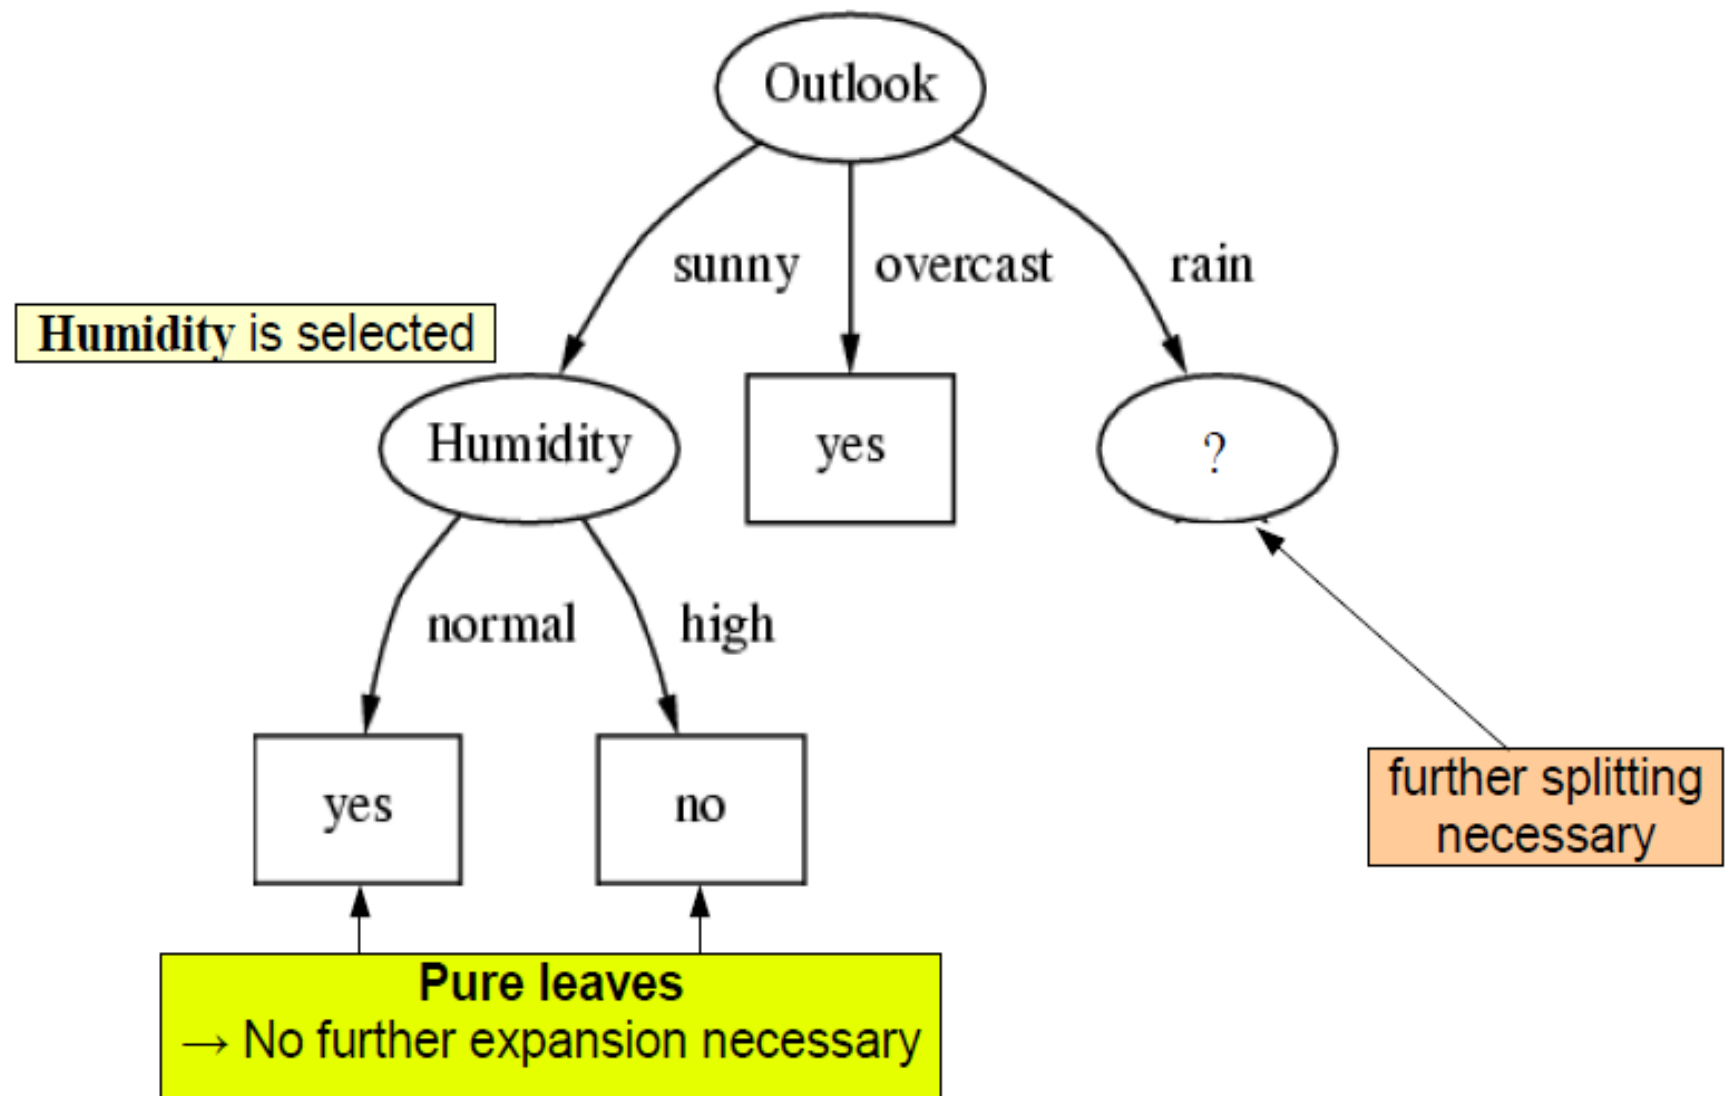


gain("Temperature") = 0.571bits

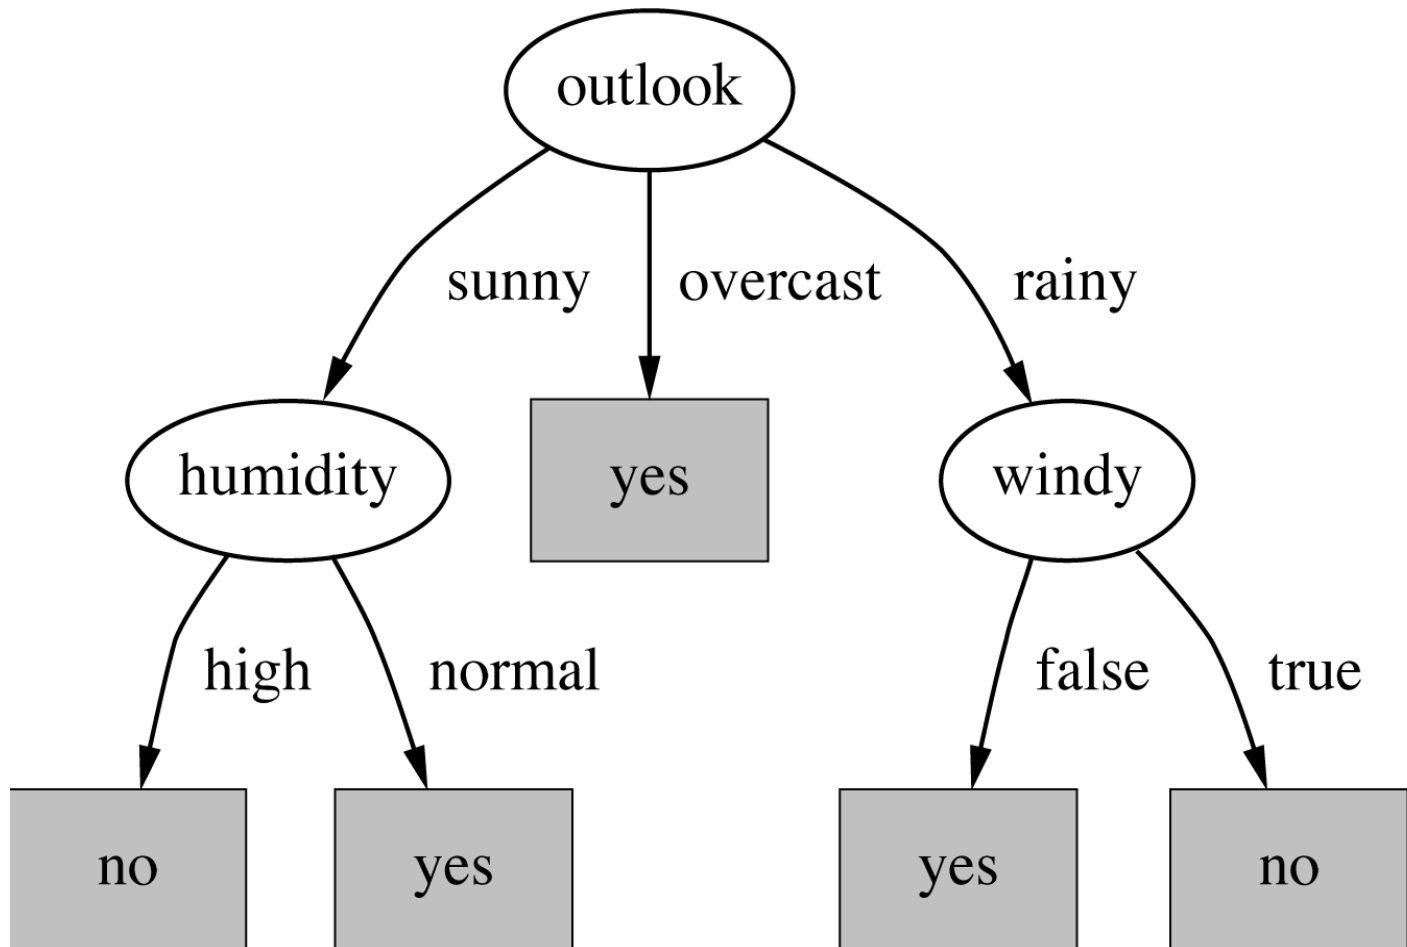
gain("Humidity") = 0.971bits

gain("Windy") = 0.020bits

Humidity is selected



The final decision tree




Highly-branching attributes

- Problematic: attributes with a large number of values (extreme case: ID code)
- Subsets are more likely to be pure if there is a large number of values
 - ⇒ Information gain is biased towards choosing attributes with a large number of values

Weather Data with ID code

ID	Outlook	Temperature	Humidity	Windy	Play?
A	sunny	hot	high	false	No
B	sunny	hot	high	true	No
C	overcast	hot	high	false	Yes
D	rain	mild	high	false	Yes
E	rain	cool	normal	false	Yes
F	rain	cool	normal	true	No
G	overcast	cool	normal	true	Yes
H	sunny	mild	high	false	No
I	sunny	cool	normal	false	Yes
J	rain	mild	normal	false	Yes
K	sunny	mild	normal	true	Yes
L	overcast	mild	high	true	Yes
M	overcast	hot	normal	false	Yes
N	rain	mild	high	true	No

Split for ID Code Attribute



Entropy of split = 0 (since each leaf node is “pure”, having only one case).

Information gain is maximal for ID code

Gain ratio

- *Gain ratio*: a modification of the information gain that reduces its bias on high-branch attributes
- Gain ratio should be
 - Large when data is evenly spread
 - Small when all data belong to one branch
- Gain ratio takes number and size of branches into account when choosing an attribute
 - It corrects the information gain by taking the *intrinsic information* of a split into account (i.e. how much info do we need to tell which branch an instance belongs to)

Gain Ratio and Intrinsic Info.

- Intrinsic information: entropy of distribution of instances into branches

$$\text{IntrinsicInfo}(S, A) \equiv -\sum \frac{|S_i|}{|S|} \log_2 \frac{|S_i|}{|S|}.$$

- *Gain ratio* (Quinlan'86) normalizes info gain by:

$$\text{GainRatio}(S, A) = \frac{\text{Gain}(S, A)}{\text{IntrinsicInfo}(S, A)}.$$

Computing the gain ratio

- Example: intrinsic information for ID code

$$\text{info}([1, 1, \dots, 1]) = 14 \times (-1/14 \times \log 1/14) = 3.807 \text{ bits}$$

- Importance of attribute decreases as intrinsic information gets larger
- Example of gain ratio:

$$\text{gain_ratio}(\text{"Attribute"}) = \frac{\text{gain}(\text{"Attribute"})}{\text{intrinsic_info}(\text{"Attribute"})}$$

- Example:

$$\text{gain_ratio}(\text{"ID_code"}) = \frac{0.940 \text{ bits}}{3.807 \text{ bits}} = 0.246$$

Gain ratios for weather data

Outlook	Temperature
Info: 0.693	Info: 0.911
Gain: $0.940 - 0.693$ 0.247	Gain: $0.940 - 0.911$ 0.029
Split info: 1.577 info([5,4,5])	Split info: 1.362 info([4,6,4])
Gain ratio: 0.156 $0.247 / 1.577$	Gain ratio: 0.021 $0.029 / 1.362$

Humidity	Windy
Info: 0.788	Info: 0.892
Gain: $0.940 - 0.788$ 0.152	Gain: $0.940 - 0.892$ 0.048
Split info: 1.000 info([7,7])	Split info: 0.985 info([8,6])
Gain ratio: $0.152 / 1$ 0.152	Gain ratio: 0.049 $0.048 / 0.985$

More on the gain ratio

- “Outlook” still comes out top
- However: “ID code” has greater gain ratio
 - Standard fix: *ad hoc* test to prevent splitting on that type of attribute
- Problem with gain ratio: it may overcompensate
 - May choose an attribute just because its intrinsic information is very low
 - Standard fix:
 - First, only consider attributes with greater than average information gain
 - Then, compare them on gain ratio

CART Splitting Criteria: Gini Index

- The Gini Index (used in CART) measures the impurity of a data partition **D**
- If a data set T contains examples from n classes, gini index, $\text{gini}(T)$ is defined as

$$\text{gini}(T) = 1 - \sum_{j=1}^n p_j^2$$

where p_j is the relative frequency of class j in T .

Gini Index

- The Gini Index considers a **binary split** for each attribute **A** say T1 and T2.

After splitting T into two subsets T1 and T2 with sizes N1 and N2, the gini index of the split data is defined as

$$gini_{split}(T) = \frac{N_1}{N} gini(T_1) + \frac{N_2}{N} gini(T_2)$$

OR

$$gini_{split}(D) = \frac{N_1}{N} gini(D1) + \frac{N_2}{N} gini(D2)$$

- The attribute providing smallest $gini_{split}(T)$ is chosen to split the node.

Gini Index

- The reduction in impurity is given by

$$\Delta Gini(A) = gini(T) - gini_{split}(T)$$

- The attribute that maximizes the reduction in impurity is chosen as the splitting attribute

Weather Data with ID code

ID	Outlook	Temperature	Humidity	Windy	Play?
A	sunny	hot	high	false	No
B	sunny	hot	high	true	No
C	overcast	hot	high	false	Yes
D	rain	mild	high	false	Yes
E	rain	cool	normal	false	Yes
F	rain	cool	normal	true	No
G	overcast	cool	normal	true	Yes
H	sunny	mild	high	false	No
I	sunny	cool	normal	false	Yes
J	rain	mild	normal	false	Yes
K	sunny	mild	normal	true	Yes
L	overcast	mild	high	true	Yes
M	overcast	hot	normal	false	Yes
N	rain	mild	high	true	No

Gini Index

- **Compute the Gini index of the training set D:** 9 tuples in class yes and 5 in class no

$$Gini(D) = 1 - \left(\left(\frac{9}{14} \right)^2 + \left(\frac{5}{14} \right)^2 \right) = 0.459$$

- **Using attribute temp:** there are three values: **cool**, **mild** and **hot**
- **Choosing the subset {cool, mild} results in two partitions:**
- **D1 (temperature ∈ {cool, mild}):** 10 tuples
- **D2 (income ∈ {hot}):** 4 tuples

Gini Index

$$\begin{aligned}
 \text{Gini}_{\text{temperature} \in \{\text{cool}, \text{mild}\} \text{ and } \{\text{hot}\}} (T) &= \frac{10}{14} \text{Gini}(D_1) + \frac{4}{14} \text{Gini}(D_2) \\
 &= \frac{10}{14} \left(1 - \left(\frac{6}{10} \right)^2 - \left(\frac{4}{10} \right)^2 \right) + \frac{4}{14} \left(1 - \left(\frac{1}{4} \right)^2 - \left(\frac{3}{4} \right)^2 \right) \\
 &= 0.450
 \end{aligned}$$

- The Gini Index measures of the remaining partitions are:
- $\text{Gini}_{\text{temperature} \in \{\text{cool}, \text{hot}\} \text{ and } \{\text{mild}\}} (T) = 0.300$
- $\text{Gini}_{\text{temperature} \in \{\text{mild}, \text{hot}\} \text{ and } \{\text{cool}\}} (T) = 0.315$
- Therefore, the best binary split for attribute temp is on { mild,hot } and {cool}

$$\Delta Gini(A) = gini(T) - gini_{split}(T)$$

$$Gini_{\text{temperature} \in \{\text{cool}, \text{mild}\} \text{ and } \{\text{hot}\}} = 0.459 - 0.450$$

Comparing Attribute Selection Measures

- The three measures, in general, return good results but
 - **Information gain:**
 - biased towards multivalued attributes
 - **Gini index:**
 - biased to multivalued attributes
 - has difficulty when # of classes is large

Bayesian Classification

Bayesian Classification

- The Bayesian Classification represents a supervised learning method as well as a statistical method for classification.
- It allows us to capture uncertainty about the model in a principled way by determining probabilities of the outcomes.
- It can solve diagnostic and predictive problems.
- This Classification is named after Thomas Bayes (1702-1761), who proposed the Bayes Theorem.
- Bayesian classification provides practical learning algorithms and prior knowledge and observed data can be combined.

Bayes' Theorem

- Let \mathbf{X} be a data tuple.
- In Bayesian terms, \mathbf{X} is considered “evidence.”
- Let H be some hypothesis, such as the data tuple \mathbf{X} belongs to a specified class C .
- For classification problems, we want to determine $P(H | \mathbf{X})$, the probability that the hypothesis H holds given the “evidence” or observed data tuple \mathbf{X}

Bayes' Theorem

$$P(H|X) = \frac{P(X|H)P(H)}{P(X)}$$

$P(H|\mathbf{X})$ is the posterior probability, or a posteriori probability, of H conditioned on \mathbf{X} .

$P(H)$ is the prior probability, or a priori probability, of H .

$P(\mathbf{X}|H)$ is the posterior probability of \mathbf{X} conditioned on H .

Naïve Bayesian Classification

- Let D be a training set of tuples and their associated class labels. As usual, each tuple is represented by an n -dimensional attribute vector, $\mathbf{X} = (x_1, x_2, \dots, x_n)$, depicting n measurements made on the tuple from n attributes, respectively, A_1, A_2, \dots, A_n .
- Suppose that there are m classes, C_1, C_2, \dots, C_m . Given a tuple, \mathbf{X} , the classifier will predict that \mathbf{X} belongs to the class having the highest posterior probability, conditioned on \mathbf{X} . That is, the naïve Bayesian classifier predicts that tuple \mathbf{X} belongs to the class C_i if and only if

$$P(C_i|\mathbf{X}) > P(C_j|\mathbf{X})$$

assign to sample \mathbf{X} the class label C such that

$P(C|\mathbf{X})$ is maximal

Naïve Bayesian Classification

$$P(C_i|X) = \frac{P(X|C_i)P(C_i)}{P(X)}.$$

- $P(X)$ is constant for all classes
- $P(C)$ = relative freq of class C samples
- C such that $P(C|X)$ is maximum =
C such that $P(X|C) \cdot P(C)$ is maximum

$$P(C_i|X) = P(X|C_i)P(C_i)$$

Weather	Play
Sunny	No
Overcast	Yes
Rainy	Yes
Sunny	Yes
Sunny	Yes
Overcast	Yes
Rainy	No
Rainy	No
Sunny	Yes
Rainy	Yes
Sunny	No
Overcast	Yes
Overcast	Yes
Rainy	No

Frequency Table		
Weather	No	Yes
Overcast		4
Rainy	3	2
Sunny	2	3
Grand Total	5	9

Likelihood table				
Weather	No	Yes		
Overcast		4	$\approx 4/14$	0.29
Rainy	3	2	$\approx 5/14$	0.36
Sunny	2	3	$\approx 5/14$	0.36
All	5	9		
	$\approx 5/14$	$\approx 9/14$		
	0.36	0.64		

Naïve Bayesian Classification

- Class prior probabilities is given by

$$P(C_i) = |C_{i,D}| / |D|,$$

where $|C_{i,D}|$ is the number of training tuples of class C_i in D .

- To compute $P(\mathbf{X} | C_i)$.

$$\begin{aligned} P(\mathbf{X} | C_i) &= \prod_{k=1}^n P(x_k | C_i) \\ &= P(x_1 | C_i) \times P(x_2 | C_i) \times \cdots \times P(x_n | C_i). \end{aligned}$$

- We can easily estimate the probabilities $P(x_1 | C_i)$, $P(x_2 | C_i)$, \dots , $P(x_n | C_i)$ from the training tuples

Naïve Bayesian Classification

- In order to predict the class label of \mathbf{X} , $P(\mathbf{X} | C_i)P(C_i)$ is evaluated for each class C_i . The classifier predicts that the class label of tuple \mathbf{X} is the class C_i if and only if
$$P(\mathbf{X} | C_i)P(C_i) > P(\mathbf{X} | C_j)P(C_j) \text{ for } 1 \leq j < m, j \neq i.$$
- In other words, the predicted class label is the class C_i for which $P(\mathbf{X} | C_i)P(C_i)$ is the maximum.

Summary

$$P(y|X) = \frac{P(X|y)P(y)}{P(X)}$$

$$X = (x_1, x_2, x_3, \dots, x_n)$$

$$P(y|x_1, \dots, x_n) = \frac{P(x_1|y)P(x_2|y)\dots P(x_n|y)P(y)}{P(x_1)P(x_2)\dots P(x_n)}$$

For all entries in the dataset, the denominator does not change, it remains static. Therefore, the denominator can be removed and proportionality can be injected.

$$P(y|x_1, \dots, x_n) \propto P(y) \prod_{i=1}^n P(x_i|y)$$

$$y = \operatorname{argmax}_y P(y) \prod_{i=1}^n P(x_i|y)$$

Naïve Bayesian Classification

- Numerical Illustration
- The dataset :

Example No.	Color	Type	Origin	Stolen?
1	Red	Sports	Domestic	Yes
2	Red	Sports	Domestic	No
3	Red	Sports	Domestic	Yes
4	Yellow	Sports	Domestic	No
5	Yellow	Sports	Imported	Yes
6	Yellow	SUV	Imported	No
7	Yellow	SUV	Imported	Yes
8	Yellow	SUV	Domestic	No
9	Red	SUV	Imported	No
10	Red	Sports	Imported	Yes

- Classify a Red Domestic SUV is getting stolen or not.

- The posterior probability $P(y|X)$ can be calculated
 - Create a **Frequency Table** for each attribute against the target.
 - Convert the frequency tables to **Likelihood Tables**
 - Use the Naïve Bayesian equation to calculate the posterior probability for each class.
 - The class with the highest posterior probability is the outcome of the prediction.
- ❖ Below are the Frequency and likelihood tables for all three predictors.

Frequency Table				Likelihood Table			
		Stolen?				Stolen?	
		Yes	No			P(Yes)	P(No)
Color	Red	3	2	Color	Red	3/5	2/5
	Yellow	2	3		Yellow	2/5	3/5

Frequency Table				Likelihood Table			
		Stolen?				Stolen?	
		Yes	No			P(Yes)	P(No)
Type	Sports	4	2	Type	Sports	4/5	2/5
	SUV	1	3		SUV	1/5	3/5

Frequency Table				Likelihood Table			
		Stolen?				Stolen?	
		Yes	No			P(Yes)	P(No)
Origin	Domestic	2	3	Origin	Domestic	2/5	3/5
	Imported	3	2		Imported	3/5	2/5

-
- ❖ Calculate the posterior probability $P(\text{Yes} | X)$ as

$$\begin{aligned} P(\text{Yes} | X) &= P(\text{Red} | \text{Yes}) * P(\text{SUV} | \text{Yes}) * P(\text{Domestic} | \text{Yes}) * P(\text{Yes}) \\ &= \frac{3}{5} * \frac{1}{5} * \frac{2}{5} * 1 \\ &= 0.048 \end{aligned}$$

and, $P(\text{No} | X)$:

$$\begin{aligned} P(\text{No} | X) &= P(\text{Red} | \text{No}) * P(\text{SUV} | \text{No}) * P(\text{Domestic} | \text{No}) * P(\text{No}) \\ &= \frac{2}{5} * \frac{3}{5} * \frac{3}{5} * 1 \\ &= 0.144 \end{aligned}$$

- ❖ Since $0.144 > 0.048$, Which means given the features RED SUV and Domestic, our example gets classified as 'NO' the car is not stolen.

Play example: estimating $P(x_i|C)$

Outlook	Temperature	Humidity	Windy	Class
sunny	hot	high	false	N
sunny	hot	high	true	N
overcast	hot	high	false	P
rain	mild	high	false	P
rain	cool	normal	false	P
rain	cool	normal	true	N
overcast	cool	normal	true	P
sunny	mild	high	false	N
sunny	cool	normal	false	P
rain	mild	normal	false	P
sunny	mild	normal	true	P
overcast	mild	high	true	P
overcast	hot	normal	false	P
rain	mild	high	true	N

$$P(p) = 9/14$$

$$P(n) = 5/14$$

outlook	
$P(\text{sunny} p) = 2/9$	$P(\text{sunny} n) = 3/5$
$P(\text{overcast} p) = 4/9$	$P(\text{overcast} n) = 0$
$P(\text{rain} p) = 3/9$	$P(\text{rain} n) = 2/5$
temperature	
$P(\text{hot} p) = 2/9$	$P(\text{hot} n) = 2/5$
$P(\text{mild} p) = 4/9$	$P(\text{mild} n) = 2/5$
$P(\text{cool} p) = 3/9$	$P(\text{cool} n) = 1/5$
humidity	
$P(\text{high} p) = 3/9$	$P(\text{high} n) = 4/5$
$P(\text{normal} p) = 6/9$	$P(\text{normal} n) = 2/5$
windy	
$P(\text{true} p) = 3/9$	$P(\text{true} n) = 3/5$
$P(\text{false} p) = 6/9$	$P(\text{false} n) = 2/5$

Play example: classifying X

- An unseen sample $X = \langle \text{rain, hot, high, false} \rangle$

$$\begin{aligned}P(X|C_i) &= \prod_{k=1}^n P(x_k|C_i) \\&= P(x_1|C_i) \times P(x_2|C_i) \times \cdots \times P(x_n|C_i).\end{aligned}$$

$$P(C_i|X) = P(X|C_i)P(C_i)$$

- $P(X|p) \cdot P(p) =$
 $P(\text{rain}|p) \cdot P(\text{hot}|p) \cdot P(\text{high}|p) \cdot P(\text{false}|p) \cdot P(p) =$
 $3/9 \cdot 2/9 \cdot 3/9 \cdot 6/9 \cdot 9/14 = 0.010582$
- $P(X|n) \cdot P(n) =$
 $P(\text{rain}|n) \cdot P(\text{hot}|n) \cdot P(\text{high}|n) \cdot P(\text{false}|n) \cdot P(n) =$
 $2/5 \cdot 2/5 \cdot 4/5 \cdot 2/5 \cdot 5/14 = 0.018286$
- Sample X is classified in class n (don't play)

- today = (Sunny, Hot, Normal, False)

$$P(Yes|today) = \frac{P(Sunny|Outlook=Yes)P(Hot|Temperature=Yes)P(Normal|Humidity=Yes)P(False|Wind=Yes)P(Yes)}{P(today)}$$

$$P(No|today) = \frac{P(Sunny|Outlook=No)P(Hot|Temperature=No)P(Normal|Humidity=No)P(False|Wind=No)P(No)}{P(today)}$$

Apply Bayesian Classification algo and construct classifier for following training set

#	Gender	Car Ownership	Travel Cost	Income level	Transportation
1	male	0	cheap	low	bus
2	male	1	cheap	medium	bus
3	female	1	cheap	medium	train
4	female	0	cheap	low	bus
5	male	1	cheap	medium	bus
6	male	0	standard	medium	train
7	female	1	standard	medium	train
8	female	1	expensive	high	car
9	male	2	expensive	medium	car
10	female	2	expensive	high	car

Classify following Samples to correct transportation class

Name	Gender	Car OwnerShip	Travel Cost	Income Level	Transportation
Alex	Male	1	Standard	High	
Buddy	Male	0	Cheap	Medium	
Cherry	Female	1	Cheap	High	










$$P(Yes|today) = \frac{P(SunnyOutlook|Yes)P(HotTemperature|Yes)P(NormalHumidity|Yes)P(NoWind|Yes)P(Yes)}{P(today)}$$

Alex	Male	1	Standard	High	
------	------	---	----------	------	--

$$\begin{aligned} P(\text{bus}|\mathbf{x}) &= P(\text{gender}|\text{bus}) * P(1|\text{bus}) * P(\text{standard}|\text{bus}) * P(\text{High}|\text{bus}) * P(\text{bus}) \\ &= 3/4 * 2/4 * 0/4 * 0/4 * 4/10 \\ &= 0 \end{aligned}$$

$$\begin{aligned} P(\text{bus}|\mathbf{x}) &= P(\text{gender}|\text{bus}) * P(1|\text{bus}) * P(\text{standard}|\text{bus}) * P(\text{High}|\text{bus}) * P(\text{bus}) \\ &= 3/4 * 2/4 * 1/7 * 1/7 * 4/10 \end{aligned}$$

Laplacian correction or Laplace estimator

Person	Hair Length	Weight	Age	Class
 Homer	0"	250	36	M
 Marge	10"	150	34	F
 Bart	2"	90	10	M
 Lisa	6"	78	8	F
 Maggie	4"	20	1	F
 Abe	1"	170	70	M
 Selma	8"	160	41	F
 Otto	10"	180	38	M
 Krusty	6"	200	45	M

	Comic	8"	290	38	?
---	-------	----	-----	----	---

Naïve Bayesian Classification

- If A_k is categorical, then $P(x_k | C_i)$ is the number of tuples of class C_i in D having the value x_k for A_k , divided by $|C_i, D|$, the number of tuples of class C_i in D .
- If A_k is continuous-valued, then we need to do a bit more work, but the calculation is pretty straightforward. A continuous-valued attribute is typically assumed to have a Gaussian distribution with a mean μ and standard deviation σ , defined by

$$g(x, \mu, \sigma) = \frac{1}{\sqrt{2\pi}\sigma} e^{-\frac{(x-\mu)^2}{2\sigma^2}}$$

$$P(x_k | C_i) = g(x_k, \mu_{C_i}, \sigma_{C_i}).$$

Overview of Naive Bayes

- The goal of Naive Bayes is to work out whether a new example is in a class given that it has a certain combination of attribute values. We work out the **likelihood of the example being in each class given the evidence (its attribute values)**, and take the highest likelihood as the classification.

- Bayes Rule: **E- Event has occurred**

$$P[H | E] = \frac{P[E | H].P[H]}{P[E]}$$

- **P[H]** is called the **prior probability** (of the hypothesis).

P[H|E] is called the **posterior probability** (of the hypothesis given the evidence)

Overview of Naive Bayes

For each class, k , work out:

$$P[H_k | E] = \frac{P[E | H_k] \cdot P[H_k]}{P[E]}$$

- Our Hypotheses are:
 - H_1 : 'the example is in class A'
 - H_2 : 'the example is in class B' etc.
- Our Evidence is the attribute values of a particular new example that is presented:
 - $E_1=x$: 'the example has value x for attribute A_1 '
 - $E_2=y$: 'the example has value y for attribute A_2 '
 - ...
 - $E_n=z$: 'the example has value z for attribute A_n '
 - Note that, *assuming the attributes are equally important and independent*, we estimate the **joint probability of that combination of attribute values** as:

$$P[E | H_k] = P[E_1 = x | H_k] \times P[E_2 = y | H_k] \times \dots \times P[E_n = z | H_k]$$

- The goal is then to find the hypothesis (**i.e. the class k**) for which the value of **$P[H_k|E]$** is at a maximum.

Overview of Naive Bayes

- For **categorical** variables we use simple proportions.

$$P[E_i=x|H_k] =$$

$$\frac{\text{no. of training egs in class } k \text{ having value } x \text{ for attribute } A_i}{\text{number of training examples in class } k}$$

- For **continuous** variables we assume a normal (Gaussian) distribution, and use the mean (μ) and standard deviation (σ) to compute the conditional probabilities.

$$P[E_i=x|H_k] = \frac{1}{\sqrt{2\pi}\sigma_k} e^{-\frac{x-\mu_k}{2\sigma_k^2}}$$

Prediction

- Predicts unknown values, i.e., models continuous-valued functions
- Predicting the identity of one thing based purely on the description of another, related thing
- Based on the relationship between a thing that you can know and a thing you need to predict
- (Numerical) prediction is similar to classification
 - construct a model
 - use model to predict continuous or ordered value for a given input
- Prediction is different from classification
 - Classification refers to predict categorical class label
 - Prediction models continuous-valued functions

Prediction

- Major method for prediction: regression
 - Regression analysis can be used to model the relationship between one or more *independent* or **predictor** variables and a *dependent* or **response** variable
- In data mining, the **predictor** variables are the attributes of interest describing the tuple. In general, the values of the predictor variables are known.
- The **response** variable is what we want to predict

How Does it Differ From Classification?

- Predicted values are usually continuous whereas classifications are discrete.
- Predictions are often (but not always) about the future whereas classifications are about the present.
- Classification is more concerned with the input than the output.

Prediction

- Regression analysis
 - Linear and multiple regression
 - Non-linear regression

predicts unknown or missing values, i.e., models continuous-valued functions

Linear Regression

- **Linear regression** : Straight-line regression analysis involves a response variable, y , and a single predictor variable, x . It is the simplest form of regression, and models y as a linear function of x .

$$y = w_0 + w_1 x$$

where the variance of y is assumed to be constant, w_0 (y-intercept) and w_1 (slope) are regression coefficients

Linear Regression


- Method of least squares:
- coefficients can be solved, which estimates the best-fitting straight line.
- Let D be a training set consisting of values of predictor variable, x , for some population and their associated values for response variable, y .
- The training set contains $|D|$ data points of the form $(x_1, y_1), (x_2, y_2), \dots, (x_{|D|}, y_{|D|})$.
- The regression coefficients can be estimated using this method with the following equations:

$$w_0 = \bar{y} - w_1 \bar{x} \qquad w_1 = \frac{\sum_{i=1}^{|D|} (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^{|D|} (x_i - \bar{x})^2}$$

- where \bar{x} is the mean value of $x_1, x_2, \dots, x_{|D|}$ and \bar{y} is the mean value of $y_1, y_2, \dots, y_{|D|}$.

Example

X(years of experience)	Y (salary) in thousands
3	30
8	57
9	64
13	72
3	36
6	43
11	59
21	90
1	20
16	83



X(years of experience)	Y (salary) in thousands				
		$x_i - x_{\text{mean}}$	$y - y_{\text{mean}}$		
3	30	-6.1	-25.4	154.94	37.21
8	57	-1.1	1.6	-1.76	1.21
9	64	-0.1	8.6	-0.86	0.01
13	72	3.9	16.6	64.74	15.21
3	36	-6.1	-19.4	118.34	37.21
6	43	-3.1	-12.4	38.44	9.61
11	59	1.9	3.6	6.84	3.61
21	90	11.9	34.6	411.74	141.61
1	20	-8.1	-35.4	286.74	65.61
16	83	6.9	27.6	190.44	47.61
				1269.6	358.9
			w_1		3.537476

Example

- Xmean=9.1 Ymean=55.4

$$w_0 = \bar{y} - w_1 \bar{x} \qquad w_1 = \frac{\sum_{i=1}^{|D|} (x_i - \bar{x})(y_i - \bar{y})}{\sum_{i=1}^{|D|} (x_i - \bar{x})^2}$$

$$w_1 = \frac{(3-9.1)(30-55.4) + (8-9.1)(57-55.4) + \dots + (16-9.1)(83-55.4)}{(3-9.1)^2 + (8-9.1)^2 + \dots + (16-9.1)^2} = 3.5$$

$$w_0 = 55.4 - (3.5)(9.1) = 23.6$$

Equation of line is: $Y = 23.6 + 3.5x$

if $x=10$ years then Salary=?

Linear Regression

- **Multiple linear regression**: involves more than one predictor variable
 - Training data is of the form $(\mathbf{x}_1, y_1), (\mathbf{x}_2, y_2), \dots, (\mathbf{x}_{|D|}, y_{|D|})$
 - Ex. we may have: $y = w_0 + w_1 x_1 + w_2 x_2$
 - Solvable by extension of least square method or using SAS, S-Plus.
 - Many nonlinear functions can be transformed into the above

Nonlinear Regression

- Some nonlinear models can be modeled by a polynomial function
- A polynomial regression model can be transformed into linear regression model. For example,

$$y = w_0 + w_1 x + w_2 x^2 + w_3 x^3$$

convertible to linear with new variables: $x_2 = x^2$, $x_3 = x^3$

$$y = w_0 + w_1 x + w_2 x_2 + w_3 x_3$$

- Other functions, such as power function, can also be transformed to linear model
- Some models are intractable nonlinear (e.g., sum of exponential terms)
 - possible to obtain least square estimates through extensive calculation on more complex formulae

Other Regression-Based Models

- Generalized linear model:
 - Logistic regression: models the prob. of some event occurring as a linear function of a set of predictor variables
 - Poisson regression: models the data that exhibit a Poisson distribution
- Log-linear models: (for categorical data)
 - Approximate discrete multidimensional prob. distributions
 - Also useful for data compression and smoothing
- Regression trees and model trees
 - Trees to predict continuous values rather than class labels

Accuracy and error measures

Classifier Evaluation Metrics: Confusion Matrix

Confusion Matrix:

Actual class\Predicted class	C_1	$\neg C_1$
C_1	True Positives (TP)	False Negatives (FN)
$\neg C_1$	False Positives (FP)	True Negatives (TN)

- Given m classes, an entry, $CM_{i,j}$ in a **confusion matrix** indicates # of tuples in class i that were labeled by the classifier as class j
- May have extra rows/columns to provide totals

A\P	C	$\neg C$	
C	TP	FN	P
$\neg C$	FP	TN	N
	P'	N'	All

- **True positives .TP/**: These refer to the positive tuples that were correctly labeled by the classifier. **Let TP be the number of true positives.**
- **True negatives .TN/**: These are the negative tuples that were correctly labeled by the classifier. **Let TN be the number of true negatives.**
- **False positives .FP/**: These are the negative tuples that were incorrectly labeled as positive (e.g., tuples of class *buys computer D no* for which the classifier predicted *buys computer D yes*). **Let FP be the number of false positives.**
- **False negatives .FN/**: These are the positive tuples that were mislabeled as negative (e.g., tuples of class *buys computer D yes* for which the classifier predicted *buys computer D no*). **Let FN be the number of false negatives.**

Classifier Evaluation Metrics: Confusion Matrix

Example of Confusion Matrix:

Actual class\Predicted class	Play= yes	Play = no	Total
Play= yes	6954	46	7000
Play = no	412	2588	3000
Total	7366	2634	10000

Classifier Accuracy, or recognition rate: percentage of test set tuples that are correctly classified

Error rate: $1 - accuracy$

Classifier Evaluation Metrics: Accuracy, Error Rate, Sensitivity and Specificity

A\P	C	¬C	
C	TP	FN	P
¬C	FP	TN	N
	P'	N'	All

- **Class Imbalance Problem:**

- One class may be *rare*, e.g. fraud, or HIV-positive
- Significant *majority of the negative class* and minority of the positive class

- **Sensitivity:** (True Positive recognition rate) proportion of actual positives which are predicted positive
 - **Sensitivity = TP/P**
- **Specificity:** (True Negative recognition rate) proportion of actual negative which are predicted negative
 - **Specificity = TN/N**
 - **Accuracy = $[Sensitivity * P / (P + N)] + [Specificity * N / (P + N)]$**
 - **Accuracy = $(TP + TN)/All$**
 - **Error rate = $(FP + FN)/All$**

-
- **Accuracy** : the proportion of the total number of predictions that were correct.
 - **Positive Predictive Value or Precision** : the proportion of positive cases that were correctly identified.
 - **Negative Predictive Value** : the proportion of negative cases that were correctly identified.
 - **Sensitivity or Recall** : the proportion of actual positive cases which are correctly identified.
 - **Specificity** : the proportion of actual negative cases which are correctly identified.

Classifier Evaluation Metrics:

Precision and Recall, and F-measures

- **Precision:** exactness – what % of tuples that the classifier labeled as positive are actually positive

$$precision = \frac{TP}{TP + FP}$$


- **Recall:** completeness – what % of positive tuples did the classifier label as positive?

$$recall = \frac{TP}{TP + FN}$$

A\P	C	¬C	
C	TP	FN	P
¬C	FP	TN	N
	P'	N'	All

Confusion Matrix		Target			
		Positive	Negative		
Model	Positive	a	b	<i>Positive Predictive Value</i>	$a/(a+b)$
	Negative	c	d	<i>Negative Predictive Value</i>	$d/(c+d)$
		<i>Sensitivity</i>	<i>Specificity</i>	Accuracy = $(a+d)/(a+b+c+d)$	
		$a/(a+c)$	$d/(b+d)$		

<https://towardsdatascience.com/beyond-accuracy-precision-and-recall-3da06bea9f6c>



How many selected items are relevant?

$$\text{Precision} = \frac{\text{true positives}}{\text{true positives} + \text{false positives}}$$

How many relevant items are selected?

$$\text{Recall} = \frac{\text{true positives}}{\text{true positives} + \text{false negatives}}$$

Classifier Evaluation Metrics: Example

Actual Class\Predicted class	cancer = yes	cancer = no	Total	Recognition(%)
cancer = yes	90	210	300	30.00 (<i>sensitivity</i>)
cancer = no	140	9560	9700	98.56 (<i>specificity</i>)
Total	230	9770	10000	96.40 (<i>accuracy</i>)

■ $Precision = 90/230 = 39.13\%$

$Recall = 90/300 = 30.00\%$

Evaluating Classifier Accuracy: Holdout & Cross-Validation Methods

- **Holdout method**
 - Given data is randomly partitioned into two independent sets
 - Training set (e.g., 2/3) for model construction
 - Test set (e.g., 1/3) for accuracy estimation
 - Random sampling: a variation of holdout
 - Repeat holdout k times, accuracy = avg. of the accuracies obtained
- **Cross-validation** (k -fold, where $k = 10$ is most popular)
 - Randomly partition the data into k *mutually exclusive* subsets, each approximately equal size
 - At i -th iteration, use D_i as test set and others as training set
 - Leave-one-out: k folds where $k = \#$ of tuples, for small sized data
 - Stratified cross-validation: folds are stratified so that class distribution of tuples in each fold is approx. the same as that in the initial data


Evaluating Classifier Accuracy: Bootstrap

- **Bootstrap**

- Works well with small data sets
- Samples the given training tuples uniformly *with replacement*
 - i.e., each time a tuple is selected, it is equally likely to be selected again and re-added to the training set

Model Selection: ROC Curves

- **ROC** (Receiver Operating Characteristics) curves: for visual comparison of classification models
- Originated from signal detection theory
- Shows the trade-off between the true positive rate and the false positive rate
- The area under the ROC curve is a measure of the accuracy of the model
- The closer to the diagonal line (i.e., the closer the area is to 0.5), the less accurate is the model



- Vertical axis represents the true positive rate
- Horizontal axis rep. the false positive rate
- The plot also shows a diagonal line
- A model with perfect accuracy will have an area of 1.0

perfect
performance

always
positive
classification

True positive rate

random
performance

liberal
performance

conservative
performance

worse than
random
performance

always
negative
classification

0.0

0.2

0.4

0.6

0.8

1.0


False positive rate

A

B

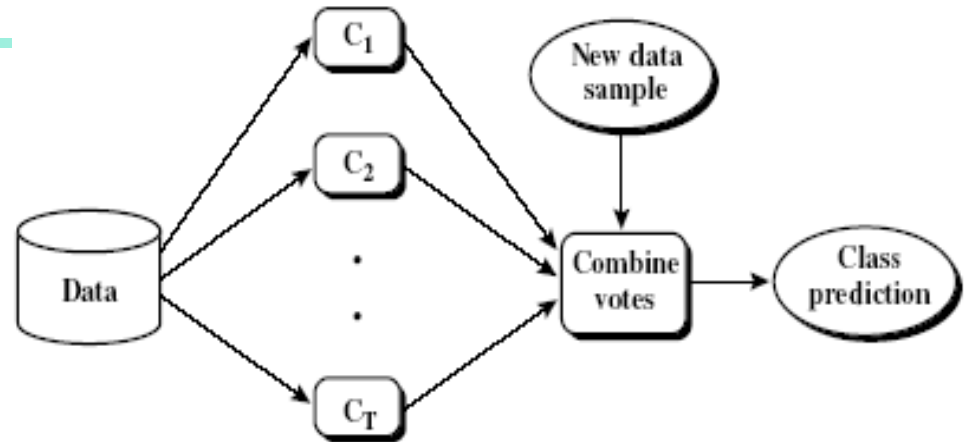
—

—



Increasing the Accuracy

Ensemble Methods: Increasing the Accuracy



- Ensemble methods
 - Use a combination of models to increase accuracy
 - Combine a series of k learned models, M_1, M_2, \dots, M_k , with the aim of creating an improved model M^*
- Popular ensemble methods
 - Bagging: averaging the prediction over a collection of classifiers
 - Boosting: weighted vote with a collection of classifiers
 - Ensemble: combining a set of heterogeneous classifiers

Bagging: Bootstrap Aggregation

- Algorithm: Bagging. The bagging algorithm—create an ensemble of models (classifiers or predictors) for a learning scheme where each model gives an equally-weighted prediction.
- Input: D , a set of d training tuples;
 - k , the number of models in the ensemble;
 - a learning scheme (e.g., decision tree algorithm etc.)
- Output: A composite model, M .

Bagging: Bootstrap Aggregation

Method:

- (1) for $i = 1$ to k do // create k models:
 - (2) create bootstrap sample, D_i , by sampling D with replacement;
 - (3) use D_i to derive a model, M_i ;
 - (4) endfor
- To use the composite model on a tuple, \mathbf{X} :
- (1) if classification then
 - (2) let each of the k models classify \mathbf{X} and return the majority vote;
 - (3) if prediction then
 - (4) let each of the k models predict a value for \mathbf{X} and return the average predicted value;

Bagging: Bootstrap Aggregation

■ Training

- Given a set D of d tuples, at each iteration i , a training set D_i of d tuples is sampled with replacement from D (i.e., bootstrap)
- A classifier model M_i is learned for each training set D_i

■ Classification: classify an unknown sample X

- Each classifier M_i returns its class prediction
- The bagged classifier M^* counts the votes and assigns the class with the most votes to X

■ Prediction: can be applied to the prediction of continuous values by taking the average value of each prediction for a given test tuple

■ Accuracy

- Often significantly better than a single classifier derived from D
- For noise data: not considerably worse, more robust
- Proved improved accuracy in prediction

Boosting

- How boosting works?
 - **Weights** are assigned to each training tuple
 - A series of k classifiers is iteratively learned
 - After a classifier M_i is learned, the weights are updated to allow the subsequent classifier, M_{i+1} , to **pay more attention to the training tuples that were misclassified** by M_i
 - The final **M^* combines the votes** of each individual classifier, where the weight of each classifier's vote is a function of its accuracy
- Boosting algorithm can be extended for numeric prediction
- Comparing with bagging: Boosting tends to have greater accuracy, but it also risks overfitting the model to misclassified data

Adaboost

- Given a set of d class-labeled tuples, $(\mathbf{X}_1, y_1), \dots, (\mathbf{X}_d, y_d)$
- Initially, all the weights of tuples are set the same ($1/d$)
- Generate k classifiers in k rounds. At round i ,
 - Tuples from D are sampled (with replacement) to form a training set D_i of the same size
 - Each tuple's chance of being selected is based on its weight
 - A classification model M_i is derived from D_i
 - Its error rate is calculated using D_i as a test set
 - If a tuple is misclassified, its weight is increased, o.w. it is decreased
- Error rate: $\text{err}(\mathbf{X}_j)$ is the misclassification error of tuple \mathbf{X}_j . Classifier M_i error rate is the sum of the weights of the misclassified tuples:

$$\text{error}(M_i) = \sum_j^d w_j \times \text{err}(\mathbf{X}_j)$$

- If $\text{error}(M_i) > 0.5$ //generate a new D_i training set, from whichn derive a new M_i .
- The weight of classifier M_i 's vote is $\log \frac{1 - \text{error}(M_i)}{\text{error}(M_i)}$

Random Forest (Breiman 2001)

- Random Forest:
 - Each classifier in the ensemble is a *decision tree* classifier and is generated using a random selection of attributes at each node to determine the split
 - During classification, each tree votes and the most popular class is returned
- Two Methods to construct Random Forest:
 - Forest-RI (*random input selection*): Randomly select, at each node, F attributes as candidates for the split at the node. The CART methodology is used to grow the trees to maximum size
 - Forest-RC (*random linear combinations*): Creates new attributes (or features) that are a linear combination of the existing attributes (reduces the correlation between individual classifiers)
- Comparable in accuracy to Adaboost, but more robust to errors and outliers
- Insensitive to the number of attributes selected for consideration at each split, and faster than bagging or boosting

Classification of Class-Imbalanced Data Sets

- Class-imbalance problem: Rare positive example but numerous negative ones, e.g., medical diagnosis, fraud, oil-spill, fault, etc.
- Traditional methods assume a balanced distribution of classes and equal error costs: not suitable for class-imbalanced data
- Typical methods for imbalance data in 2-class classification:
 - **Oversampling:** re-sampling of data from positive class
 - **Under-sampling:** randomly eliminate tuples from negative class
 - **Threshold-moving:** moves the decision threshold, t , so that the rare class tuples are easier to classify, and hence, less chance of costly false negative errors
 - Ensemble techniques: Ensemble multiple classifiers introduced above
- Still difficult for class imbalance problem on multiclass tasks

$$\begin{aligned}
 P(X|C_i) &= \prod_{k=1}^n P(x_k|C_i) \\
 &= P(x_1|C_i) \times P(x_2|C_i) \times \cdots \times P(x_n|C_i).
 \end{aligned}$$

e 1

$$P(C_i|X) = P(X|C_i)P(C_i)$$

Take the following training data, from bank loan applicants:

ApplicantID	City	Children	Income	Status
1	Delhi	Many	Medium	DEFAULTS
2	Delhi	Many	Low	DEFAULTS
3	Delhi	Few	Medium	PAYS
4	Delhi	Few	High	PAYS

- $P[\text{City}=\text{Delhi} \mid \text{Status} = \text{DEFAULTS}] = 2/2 = 1$
- $P[\text{City}=\text{Delhi} \mid \text{Status} = \text{PAYS}] = 2/2 = 1$
- $P[\text{Children}=\text{Many} \mid \text{Status} = \text{DEFAULTS}] = 2/2 = 1$
- $P[\text{Children}=\text{Few} \mid \text{Status} = \text{DEFAULTS}] = 0/2 = 0$

Worked Example 1

Summarizing, we have the following probabilities:

Probability of...	... given DEFAULTS	... given PAYS
City=Delhi	$2/2 = 1$	$2/2 = 1$
Children=Few	$0/2 = 0$	$2/2 = 1$
Children=Many	$2/2 = 1$	$0/2 = 0$
Income=Low	$1/2 = 0.5$	$0/2 = 0$
Income=Medium	$1/2 = 0.5$	$1/2 = 0.5$
Income=High	$0/2 = 0$	$1/2 = 0.5$

and $P[\text{Status} = \text{DEFAULTS}] = 2/4 = 0.5$

$P[\text{Status} = \text{PAYS}] = 2/4 = 0.5$

The probability of `Income=Medium` given the applicant `DEFAULTS` =
the number of applicants with `Income=Medium` who `DEFAULT`
divided by the number of applicants who `DEFAULT`
= $1/2 = 0.5$

Worked Example 1

Now, assume a new example is presented where

City=Delhi, Children=Many, and Income=Medium:

First, we estimate the likelihood that the example is a defaulter, given its attribute values

$$\begin{aligned} P[\text{Status} = \text{DEFAULTS} \mid \text{Delhi, Many, Medium}] &= \\ &P[\text{Delhi} \mid \text{DEFAULTS}] * \\ &P[\text{Many} \mid \text{DEFAULTS}] * \\ &P[\text{Medium} \mid \text{DEFAULTS}] * \\ &P[\text{DEFAULTS}] \\ &= 1 * 1 * 0.5 * 0.5 \\ &= 0.25 \end{aligned}$$

Then we estimate the likelihood that the example is a payer, given its attributes

$$\begin{aligned} P[\text{Status} = \text{PAYS} \mid \text{Delhi, Many, Medium}] &= \\ &P[\text{Delhi} \mid \text{PAYS}] * \\ &P[\text{Many} \mid \text{PAYS}] * \\ &P[\text{Medium} \mid \text{PAYS}] * \\ &P[\text{PAYS}] \\ &= 1 * 0 * 0.5 * 0.5 = 0 \end{aligned}$$

As the conditional likelihood of being a defaulter is higher (because $0.25 > 0$), we conclude that the new example is a defaulter.

Worked Example 1

Now, assume a new example is presented where

City=Delhi, Children=Many, and Income=High:

First, we estimate the likelihood that the example is a defaulter, given its attribute values:

$$\begin{aligned} P[\text{Status} = \text{DEFAULTS} \mid \text{Delhi}, \text{Many}, \text{High}] &= \\ P[\text{Delhi} \mid \text{DEFAULTS}] \times P[\text{Many} \mid \text{DEFAULTS}] \times P[\text{High} \mid \text{DEFAULTS}] \times \\ P[\text{DEFAULTS}] \\ &= 1 \times 1 \times 0 \times 0.5 = 0 \end{aligned}$$

Then we estimate the likelihood that the example is a payer, given its attributes:

$$\begin{aligned} P[\text{Status} = \text{PAYS} \mid \text{Delhi}, \text{Many}, \text{High}] &= \\ P[\text{Delhi} \mid \text{PAYS}] \times P[\text{Many} \mid \text{PAYS}] \times P[\text{High} \mid \text{PAYS}] \times P[\text{PAYS}] \\ &= 1 \times 0 \times 0.5 \times 0.5 = 0 \end{aligned}$$

As the conditional likelihood of being a defaulter is the same as that for being a payer, we can come to no conclusion for this example.

Worked Example 2

Take the following training data, for credit card authorizations:

TransactionID	Income	Credit	Decision
1	Very High	Excellent	AUTHORIZE
2	High	Good	AUTHORIZE
3	Medium	Excellent	AUTHORIZE
4	High	Good	AUTHORIZE
5	Very High	Good	AUTHORIZE
6	Medium	Excellent	AUTHORIZE
7	High	Bad	REQUEST ID
8	Medium	Bad	REQUEST ID
9	High	Bad	REJECT
10	Low	Bad	CALL POLICE

Assume we'd like to determine how to classify a new transaction,
with Income = Medium and Credit=Good.

Worked Example 2

Our conditional probabilities are:

Probability of...	... given AUTHORIZE	... given REQUEST ID	... given REJECT	... given CALL POLICE
Income=Very High	2/6	0/2	0/1	0/1
Income=High	2/6	1/2	1/1	0/1
Income=Medium	2/6	1/2	0/1	0/1
Income=Low	0/6	0/2	0/1	1/1
Credit=Excellent	3/6	0/2	0/1	0/1
Credit=Good	3/6	0/2	0/1	0/1
Credit=Bad	0/6	2/2	1/1	1/1

Our class probabilities are:

$$P[\text{Decision} = \text{AUTHORIZE}] = 6/10$$

$$P[\text{Decision} = \text{REQUEST ID}] = 2/10$$

$$P[\text{Decision} = \text{REJECT}] = 1/10$$

$$P[\text{Decision} = \text{CALL POLICE}] = 1/10$$

Weaknesses

- Naive Bayes assumes that variables are **equally important** and that they are **independent** which is often not the case in practice.
- Naive Bayes is **damaged by the inclusion of redundant (strongly dependent) attributes**. e.g. if people with high income have expensive houses, then including both income and house-price in the model would unfairly multiply the effect of having low income.
- Sparse data: If some attribute values are not present in the data, then a **zero probability** for $P[X|C]$ might exist. This would lead $P[C|X]$ to be zero no matter how high $P[X|C]$ is for other attribute values.